

CSE 486/586 Distributed Systems Consistency --- 3

Steve Ko
Computer Sciences and Engineering
University at Buffalo

CSE 486/586, Spring 2012

Recap

- Consistency
 - Linearizability?
 - Sequential consistency?
- Chain replication
- Primary-backup (passive) replication
- Active replication

CSE 486/586, Spring 2012

2

Linearizability vs. Sequential Consistency

- Both care about giving **an illusion of a single copy**.
 - From the outside observer, the system should (almost) behave as if there's only a single copy.
- Linearizability cares about **time**.
 - Steve writes on his facebook wall at 11am.
 - Atri writes on his facebook wall at 11:05am.
 - Everyone will see the posts in that order.
- Sequential consistency cares about **program order**.
 - Steve writes on his facebook wall at 11am.
 - Atri writes on his facebook wall at 11:05am.
 - It's not necessarily that the posts will be ordered that way (though everyone will see the same order).

CSE 486/586, Spring 2012

3

Two More Consistency Models

- Even more relaxed
 - We don't even care about providing an illusion of a single copy.
- Causal consistency
 - We care about ordering causally related write operations correctly.
- Eventual consistency
 - As long as we can say all replicas converge to the same copy eventually, we're fine.

CSE 486/586, Spring 2012

4

Causal Consistency

- Writes that are potentially causally related must be seen by all processes in the same order. Concurrent writes may be seen in a different order on different machines.
 - Weaker than sequential consistency
- How do we define "causal relations" between two writes?
 - (Roughly) One client reads something that another client has written; then the client writes something.




CSE 486/586, Spring 2012

5

Causal Consistency



- Example 1:




CSE 486/586, Spring 2012

6

Causal Consistency Example 2

- Causally consistent?




- No!

CSE 486/586, Spring 2012

7

Causal Consistency Example 3

- Causally consistent?




- Yes!

CSE 486/586, Spring 2012

8

Eventual Consistency

- Popularized by the CAP theorem.
- The main problem is network partitions.



CSE 486/586, Spring 2012

9

Dilemma

- In the presence of a **network partition**:
- In order to keep the replicas **consistent**, you need to block.
 - From the outside observer, the system appears to be **unavailable**.
- If we still serve the requests from two partitions, then the replicas will diverge.
 - The system is **available**, but no **consistency**.
- The CAP theorem explains this dilemma.

CSE 486/586, Spring 2012

10

CAP Theorem

- **Consistency**
- **Availability**
 - Respond with a reasonable delay
- **Partition tolerance**
 - Even if the network gets partitioned
- Choose two!
- Brewer conjectured in 2000, then proven by Gilbert and Lynch in 2002.

CSE 486/586, Spring 2012

11

Coping with CAP

- The main issue is the Internet.
 - As the system grows to span geographically distributed areas, network partitioning becomes inevitable.
- Then the choice is either giving up availability or consistency
- A design choice: What makes more sense to your scenario?
- Giving up availability and retaining consistency
 - E.g., use 2PC
 - Your system blocks until everything becomes consistent.
- Giving up consistency and retaining availability
 - Eventual consistency

CSE 486/586, Spring 2012

12

CSE 486/586 Administrivia

- PA4 will be released soon.
- Anonymous feedback form still available.
- Please come talk to me!

CSE 486/586, Spring 2012

13

Dealing with Network Partitions

- During a partition, pairs of conflicting transactions may have been allowed to execute in different partitions. The only choice is to take corrective action after the network has recovered
 - Assumption: Partitions heal eventually
- Abort one of the transactions after the partition has healed
- Basic idea: allow operations to continue in one or some of the partitions, but reconcile the differences later after partitions have healed

CSE 486/586, Spring 2012

14

Quorum Approaches

- Quorum approaches used to decide whether reads and writes are allowed
- There are two types: pessimistic quorums and optimistic quorums
- In the pessimistic quorum philosophy, updates are allowed only in a partition that has the majority of RMs
 - Updates are then propagated to the other RMs when the partition is repaired.

CSE 486/586, Spring 2012

15



Static Quorums

- The decision about how many RMs should be involved in an operation on replicated data is called Quorum selection
- Quorum rules state that:
 - At least r replicas must be accessed for read
 - At least w replicas must be accessed for write
 - $r + w > N$, where N is the number of replicas
 - $w > N/2$
 - Each object has a version number or a consistent timestamp

CSE 486/586, Spring 2012

16

Static Quorums

-  • What does $r + w > N$ mean?
 - The only way to satisfy this condition is that there's always an overlap between the reader set and the write set.
 - There's always some replica that has the most recent write.
-  • What does $w > N/2$ mean?
 - When there's a network partition, only the partition with more than half of the RMs can perform write operations.
 - The rest will just serve reads with stale data.
- R and W are tunable:
 - E.g., $N=3, r=1, w=3$: High read throughput, perhaps at the cost of write throughput.

CSE 486/586, Spring 2012

17

Optimistic Quorum Approaches

- An Optimistic Quorum selection allows writes to proceed in any partition.
- "Write, but don't commit"
 - Unless the partition gets healed in time.
- Resolve write-write conflicts after the partition heals.
- Optimistic Quorum is practical when:
 - Conflicting updates are rare
 - Conflicts are always detectable
 - Damage from conflicts can be easily confined
 - Repair of damaged data is possible or an update can be discarded without consequences
 - Partitions are relatively short-lived

CSE 486/586, Spring 2012

18

View-based Quorum

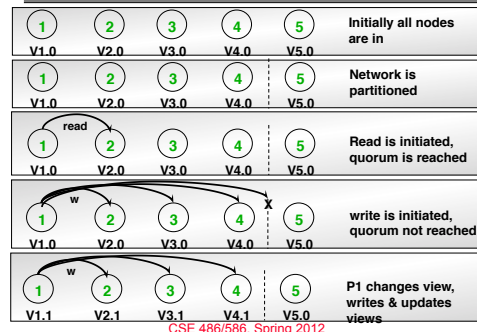
- An optimistic approach
- Quorum is based on views at any time
 - Uses group communication as a building block (see previous lecture)
- We define thresholds for each of read and write :
 - W : regular writer quorum
 - R : regular reader quorum
 - A_w : minimum nodes in a view for write, e.g., $A_w > N/4$
 - A_r : minimum nodes in a view for read
 - E.g., $A_w + A_r > N/2$
- Protocol
 - Try regular quorum first; if it doesn't work, change the view. If the minimum is satisfied, then proceed.
 - A_w & A_r effectively determine which partition can proceed.

CSE 486/586, Spring 2012

19

Example: View-based Quorum

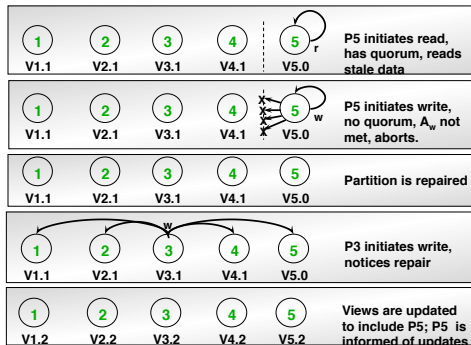
- Consider: $N = 5$, $w = 5$, $r = 1$, $A_w = 3$, $A_r = 1$



CSE 486/586, Spring 2012

20

Example: View-based Quorum (cont'd)



CSE 486/586, Spring 2012

21

Summary

- Causal consistency & eventual consistency
- Quorums
 - Static
 - Optimistic
 - View-based

CSE 486/586, Spring 2012

22

Acknowledgements

- These slides contain material developed and copyrighted by Indranil Gupta (UIUC).

CSE 486/586, Spring 2012

23