

Outline

❑ *Concepts, classical CPU virtualization*

- *Basic interpretation*

❑ *Memory virtualization*

Memory allocation

- ❑ Each VM usually receives a contiguous set of physical addresses.
 - 512 Mbyte – 4 Gbyte are typical values.
- ❑ As far as VM is concerned, this is the physical memory of the machine.
- ❑ The guest OS allocates pages or segments to guest processes.

Memory management

❑ Assumptions of OS in VM:

- Physical memory is a contiguous block of addresses from 0 to some n.
- OS can map any virtual page to any page frame.

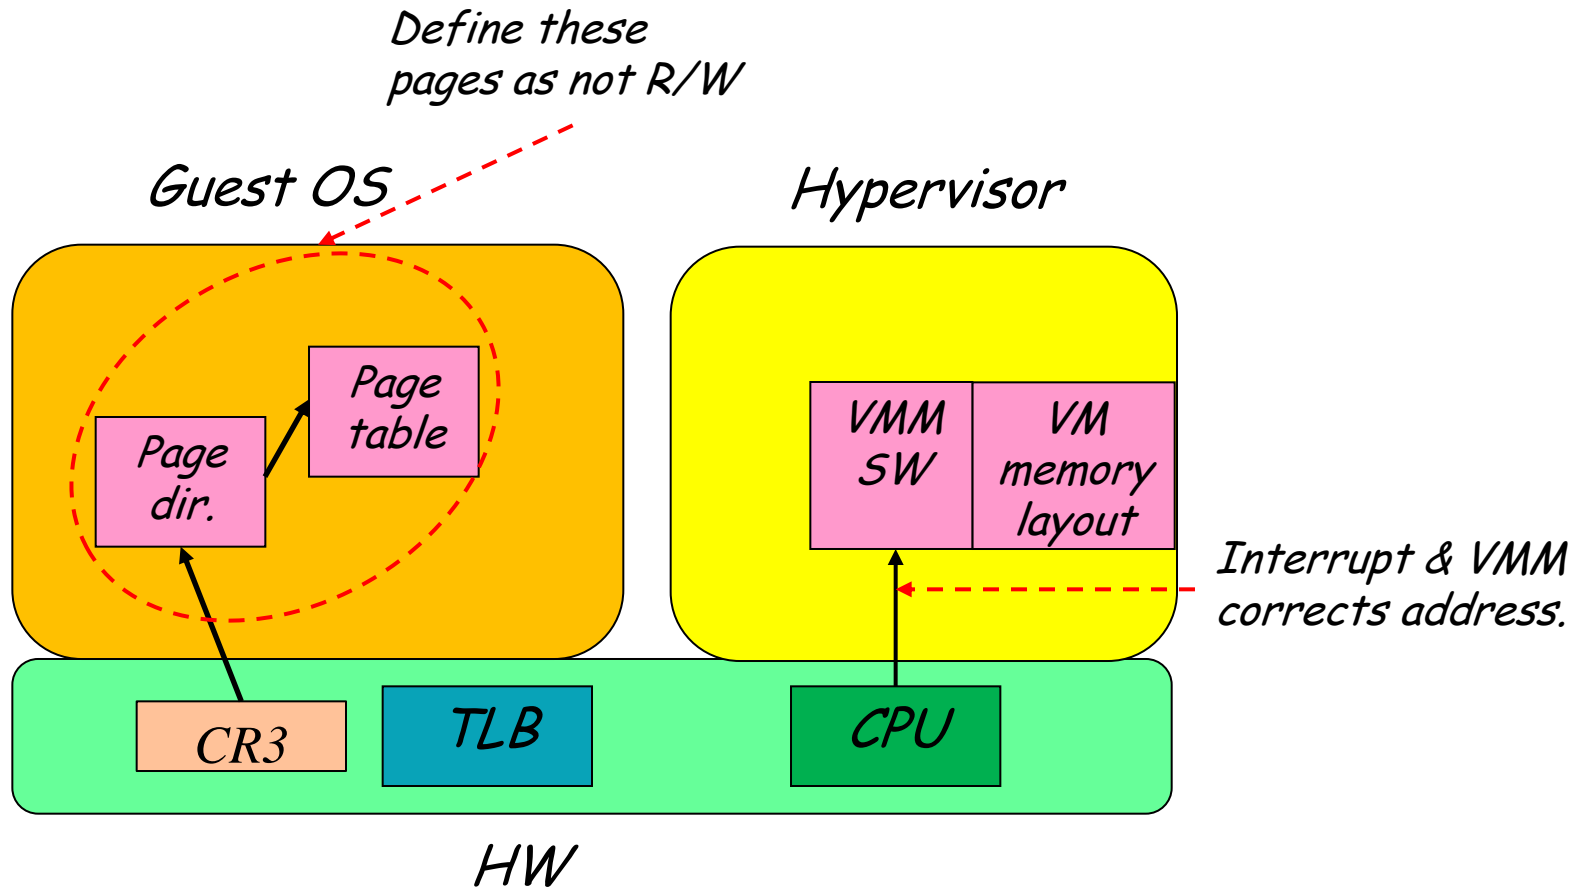
❑ Hypervisor must:

- Partition memory among VMs.
- Ensure virtual page mapping only to assigned page frames.

❑ TLB – page fault in HW-managed TLB (e.g. x86) causes HW to select a page from page table.

❑ VM OS must not manage real page table.

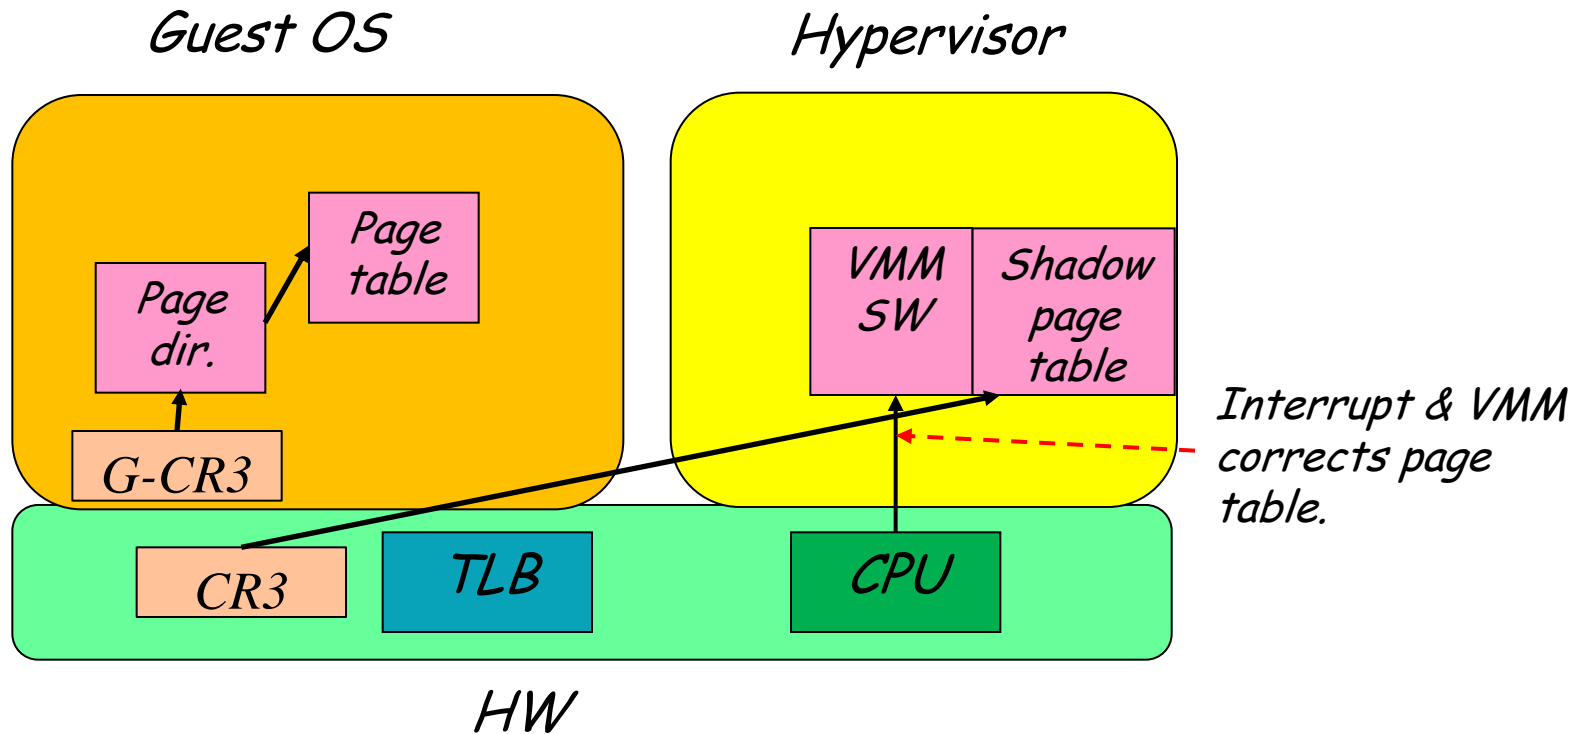
Option 1: brute force



Brute force – description

- ❑ Guest page tables are read and write protected in host system.
- ❑ If guest OS reads page table (e.g. for page eviction) writes page table (e.g. after page fault), or changes CR3, the system traps.
- ❑ The hypervisor then uses a VM memory layout to:
 - ❑ Return answers to VM
 - ❑ Update the layout
- ❑ Hypervisor switches VM memory layout when new VM is scheduled.

Option 2: shadow page tables



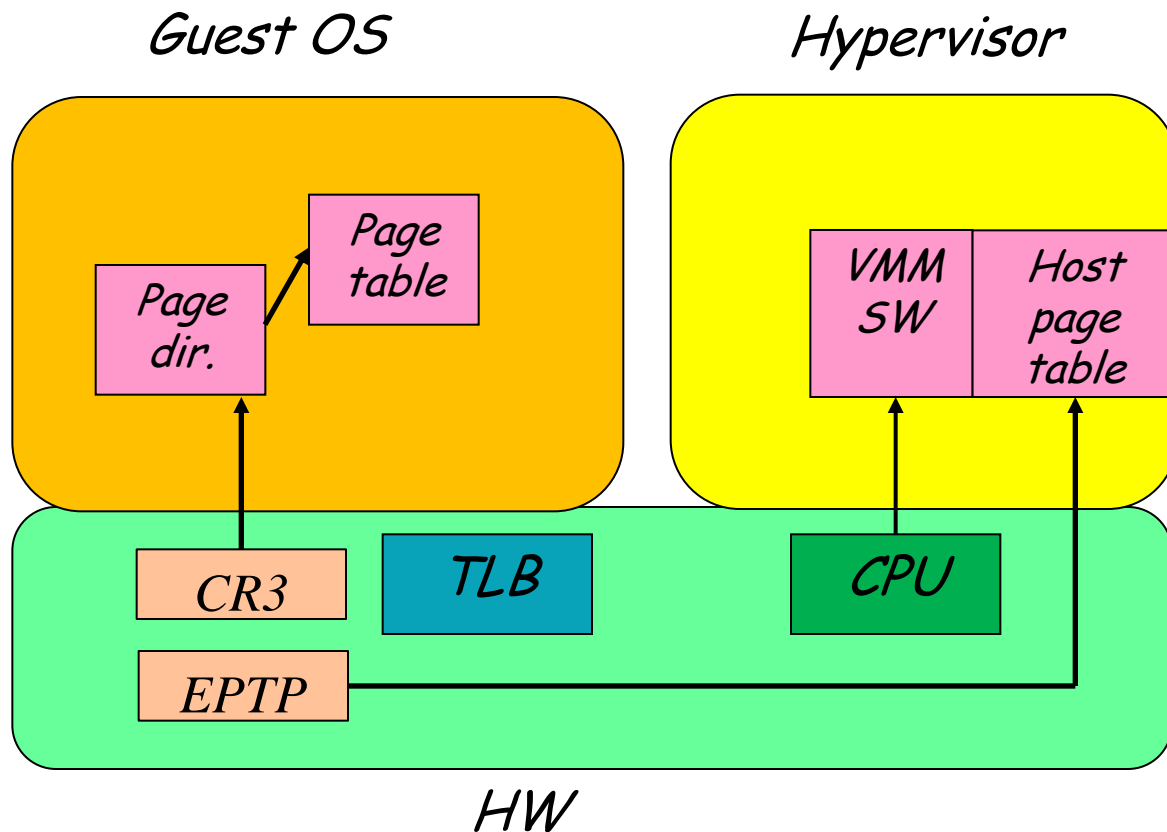
Shadow page tables – description

- ❑ Hypervisor maintains “shadow page tables”.
- ❑ Guest page tables map: Guest VA → Guest PA
- ❑ Shadow tables: Guest VA → Host PA.
- ❑ Hypervisor does not trap guest updates to its page table.
 - Result – inconsistent guest page table and shadow page table.
- ❑ When guest process accesses virtual address
 - The physical address is not in the guest page table, but in the shadow page table.
 - HW translates correctly, because it is aware only of shadow tables.

Shadow page tables – description (continued)

- ❑ If address in TLB – TLB hit and no problem.
- ❑ When guest process causes a page fault
 - Hypervisor begins execution.
 - Hypervisor updates guest page table with new page.
 - Hypervisor updates shadow page table.
- ❑ Performance is as good as native execution as long as there are no page faults.
- ❑ Shadow page tables should be cached so that once a VM is re-scheduled the page table does not have to be rebuilt from scratch.

Option 3: nested page tables



Nested page tables - description

- ❑ The name implies having page tables within page tables.
- ❑ The essence of the idea is a **hardware assist**.
 - Hardware has an extra pointer and the ability to walk an extra set of page tables.
 - Idea is called Extended Page Tables (EPT) by Intel
- ❑ Guest page tables hold Guest VA → Guest PA mapping, access by standard CR3
- ❑ Extended page tables hold Host VA → Host PA mapping, access by EPTP (EPT pointer).
- ❑ Host VA=Guest PA

Nested page tables – description (cont'd)

❑ TLB as usual holds Guest VA → Host PA

❑ On memory access

- If found in TLB – no problem.
- If not in TLB, but no page fault, hardware walks **both tables** and updates TLB.
- If page fault, then hardware hypervisor gets host physical page and provides host virtual page (guest physical) to VM.

Sources

- ❑ “Modern operating systems”, 4th edition, A. Tanenbaum and H. Bos
- ❑ “Virtual machines”, J. E. Smith and R. Nair
- ❑ A presentation by Niv Gilboa from CSE@BGU
- ❑ “Formal requirements for virtualizable third generation architectures”, G. J. Popek and R. P. Goldberg, CACM, 1974
- ❑ “A comparison of software and hardware techniques for x86 virtualization”, K. Adams and O. Ageson, ASPLOS 2006