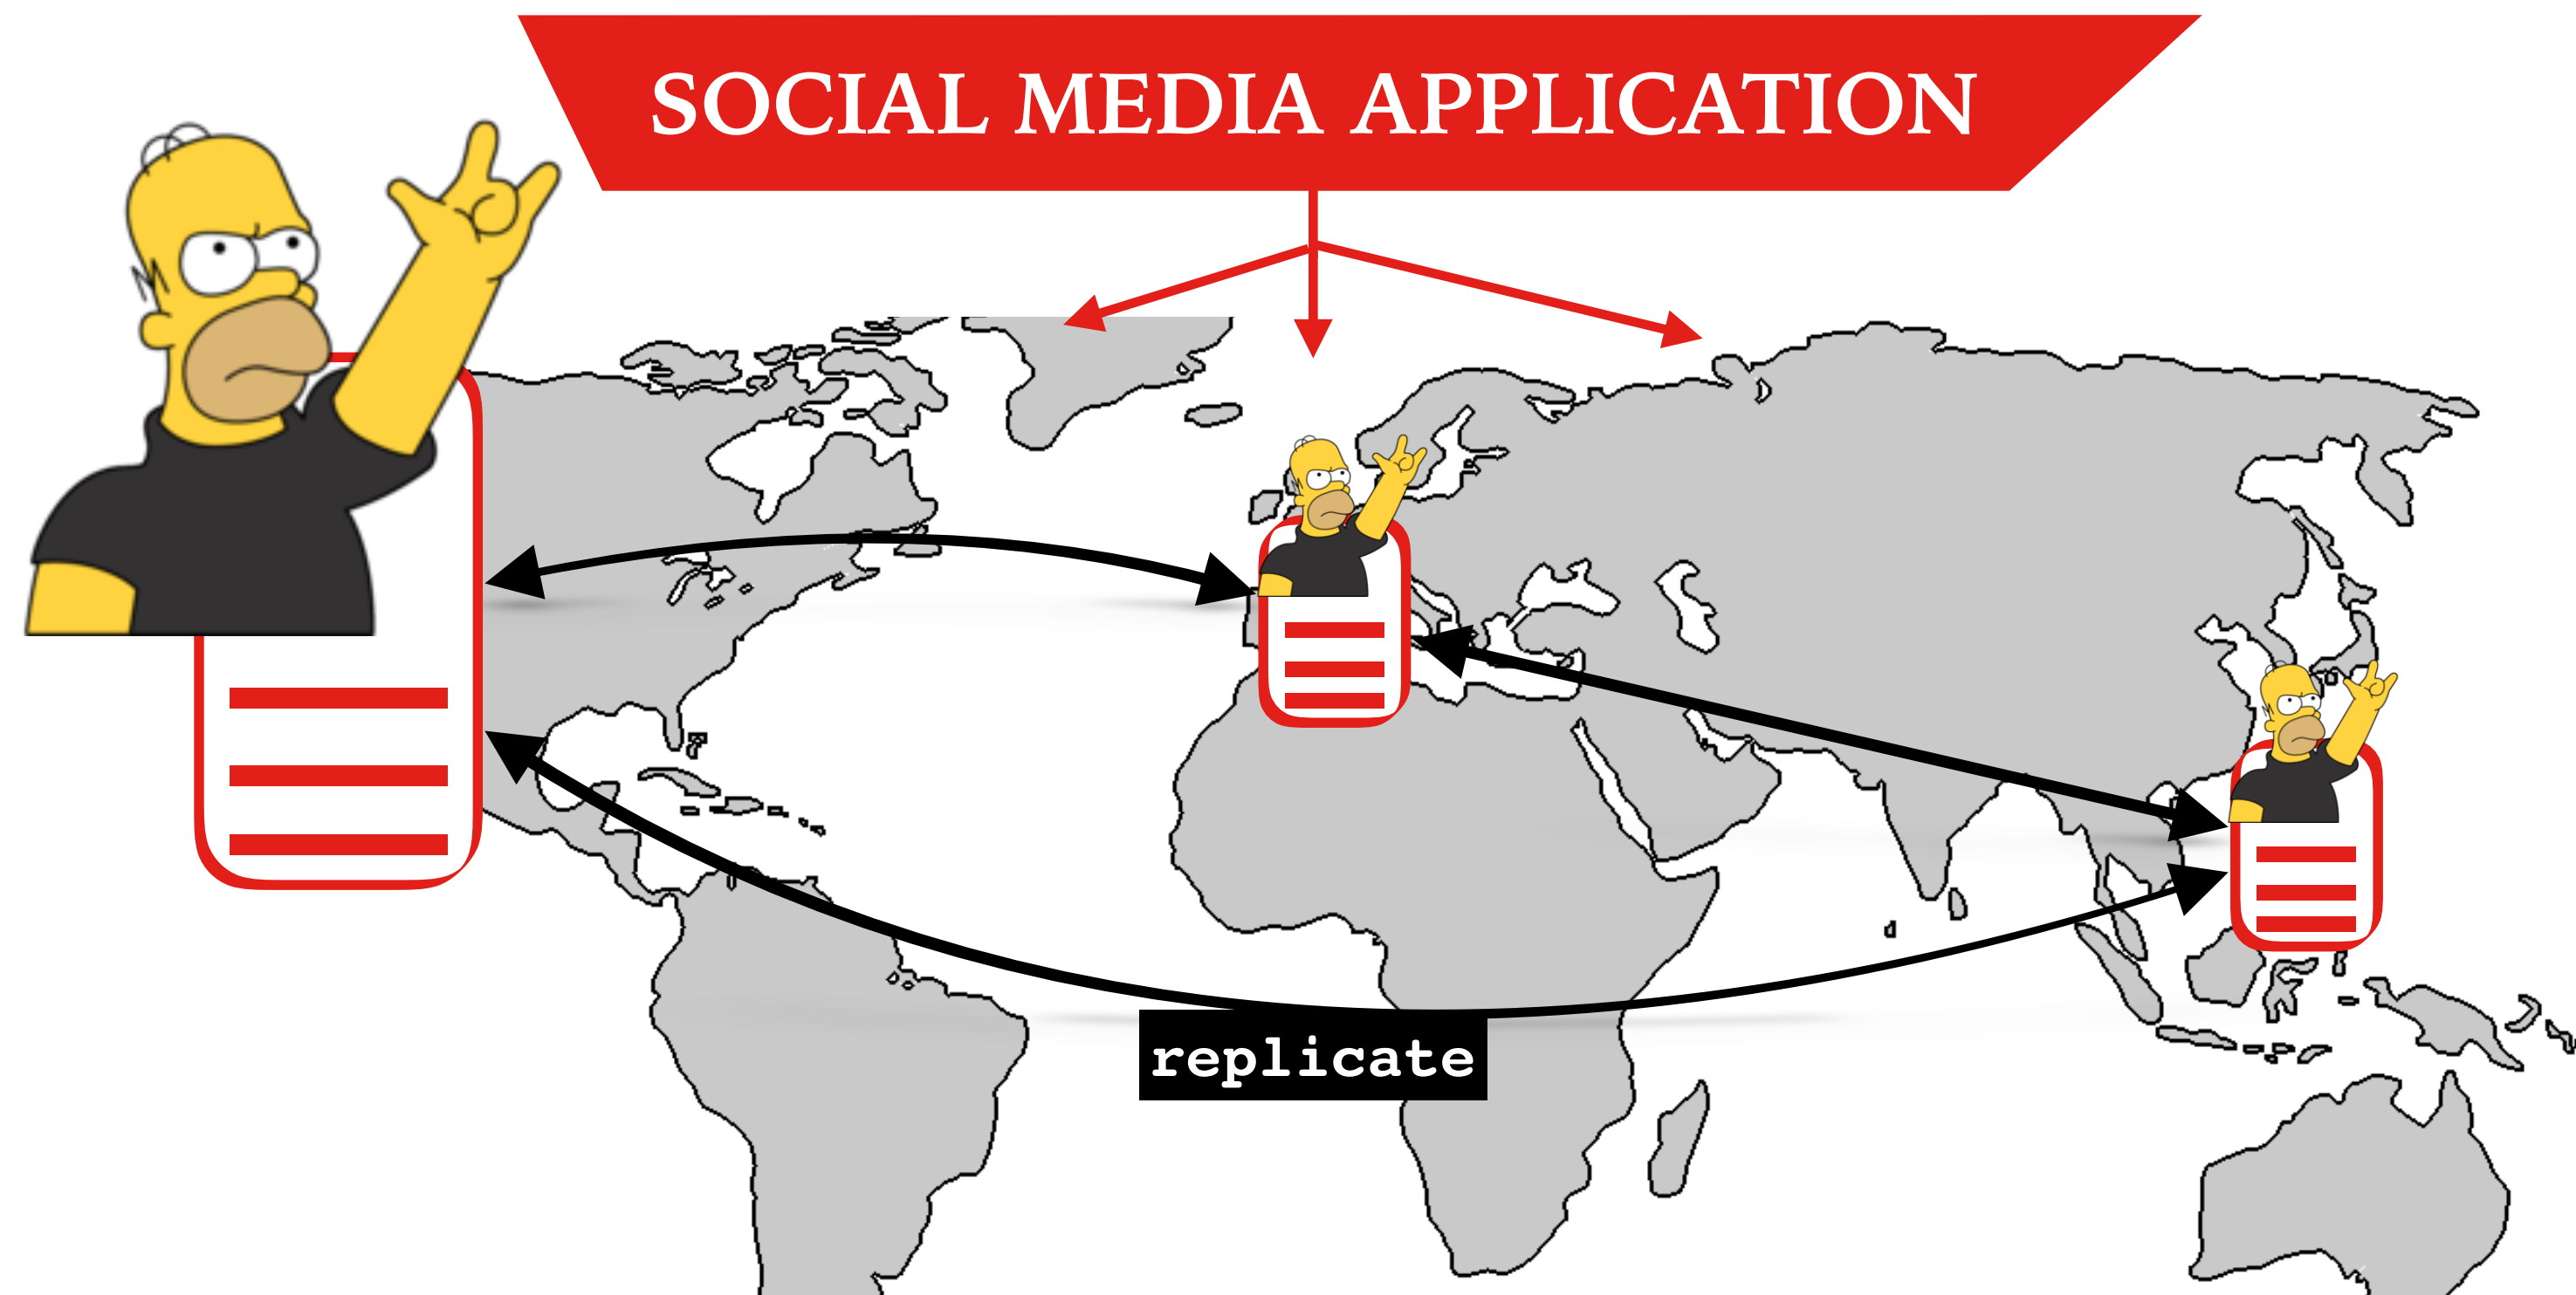
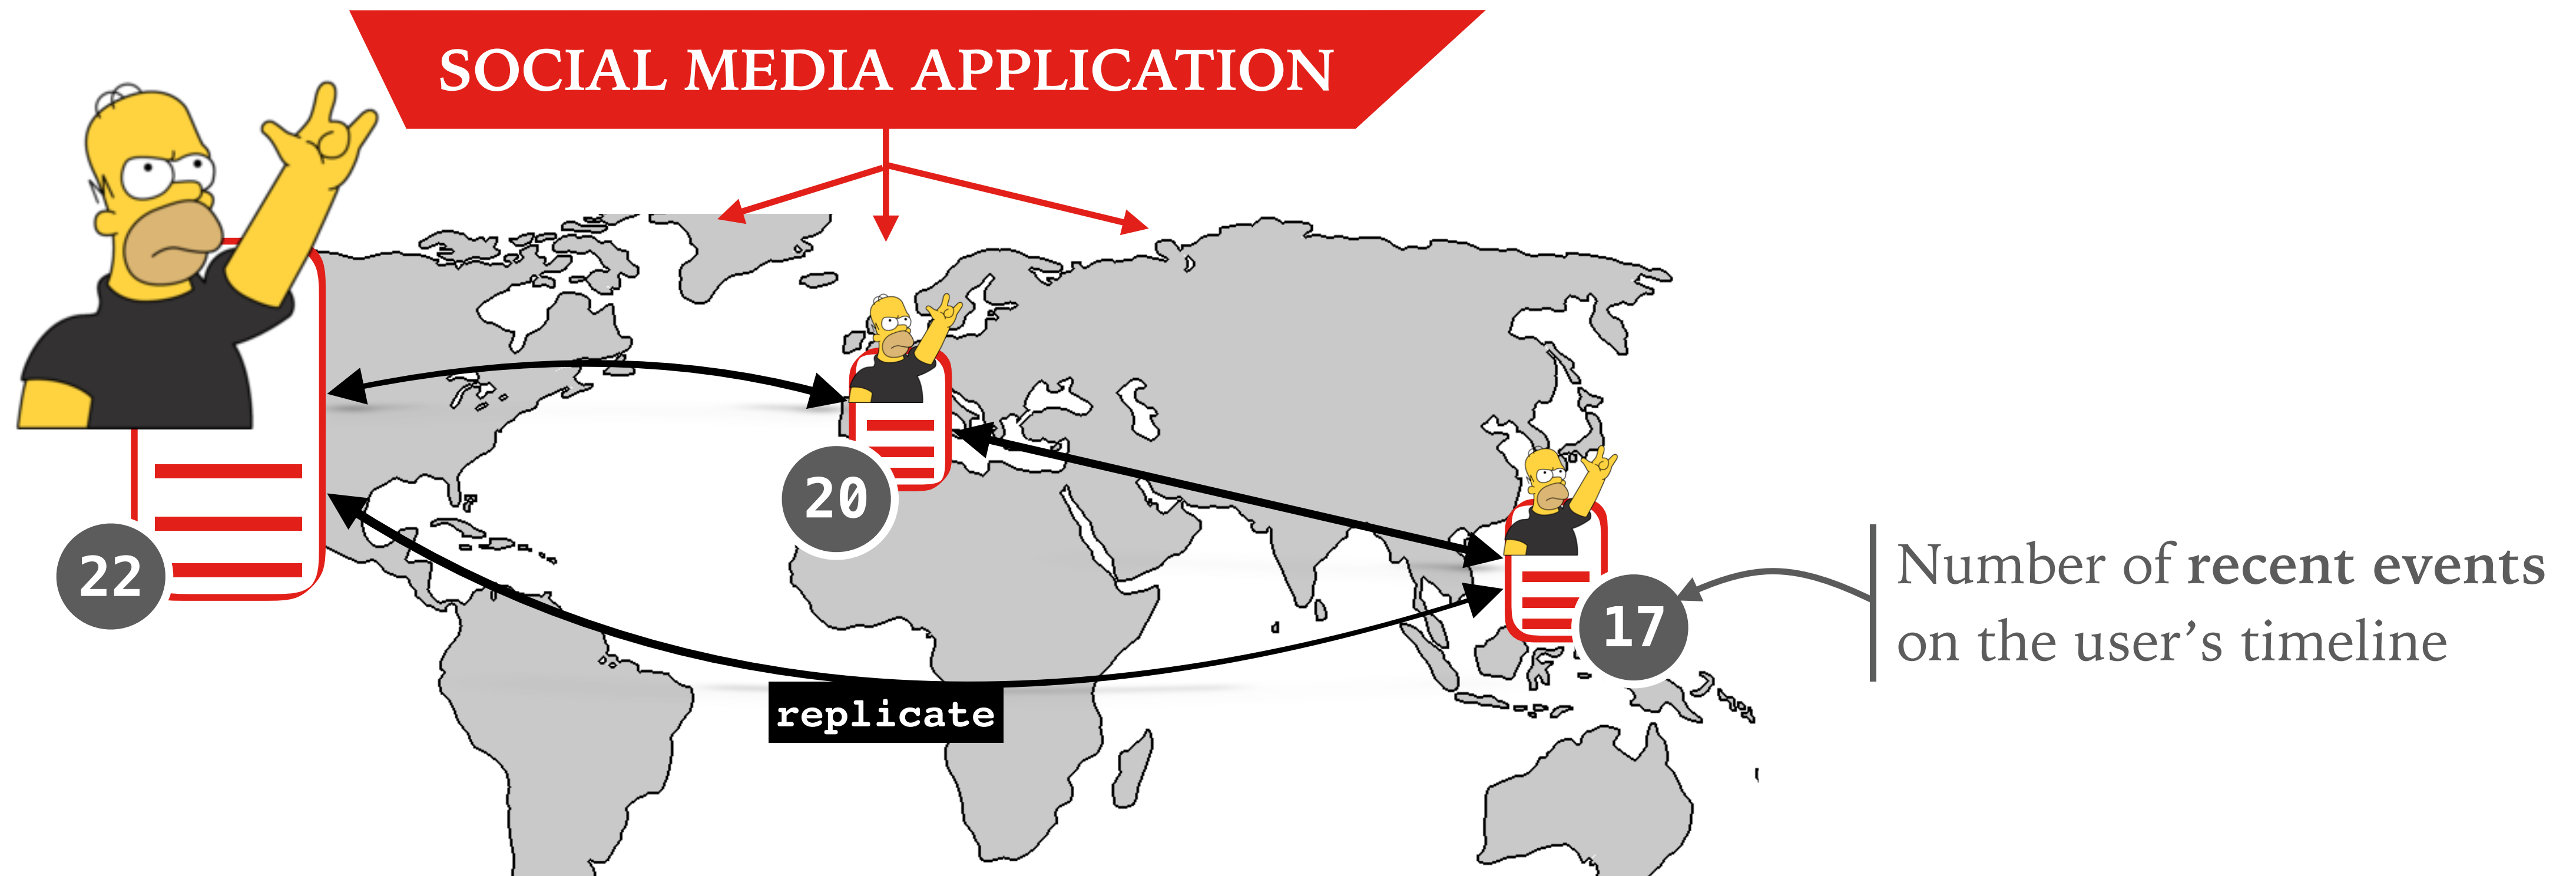


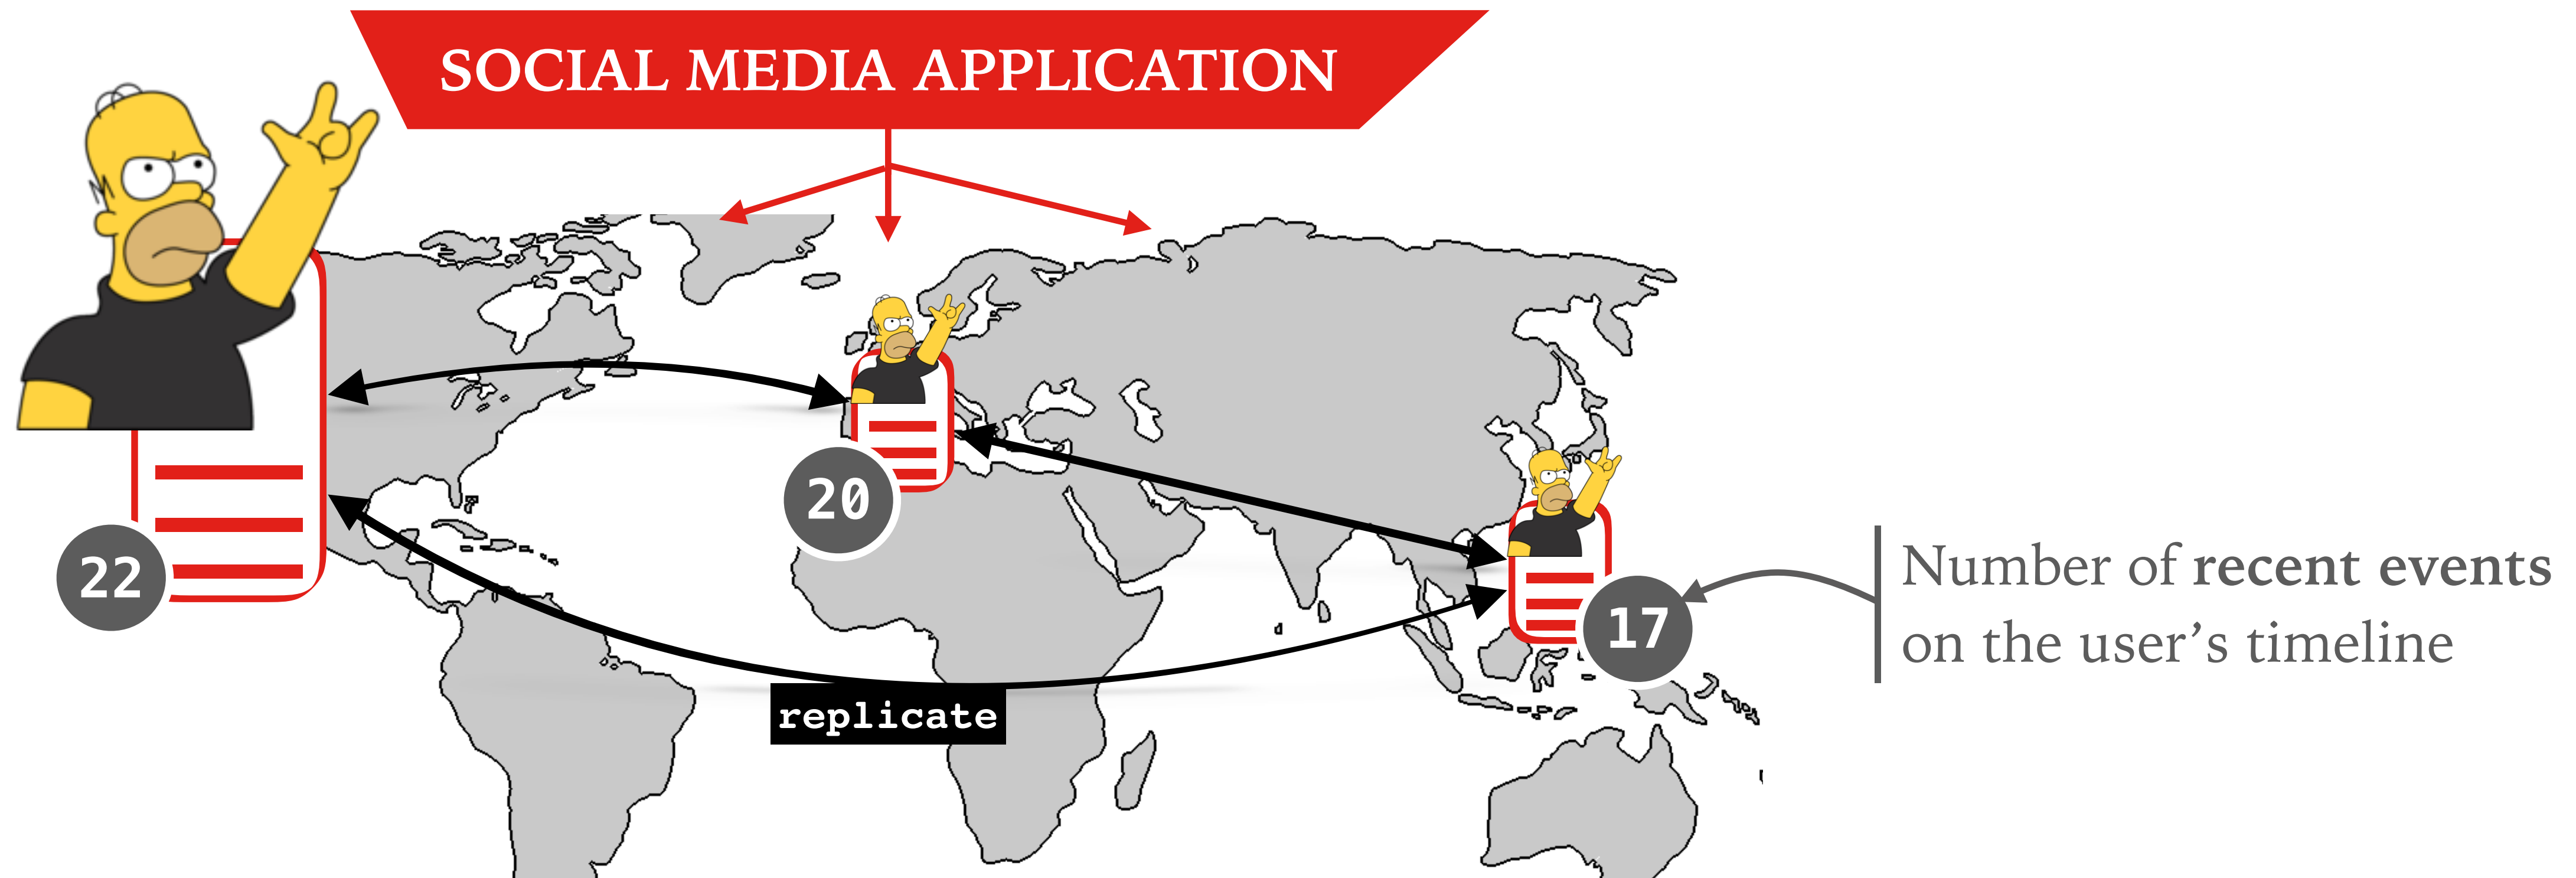
Incremental Consistency Guarantees

For Replicated Objects

Rachid Guerraoui, Matej Pavlovic, Dragos-Adrian Seredinschi

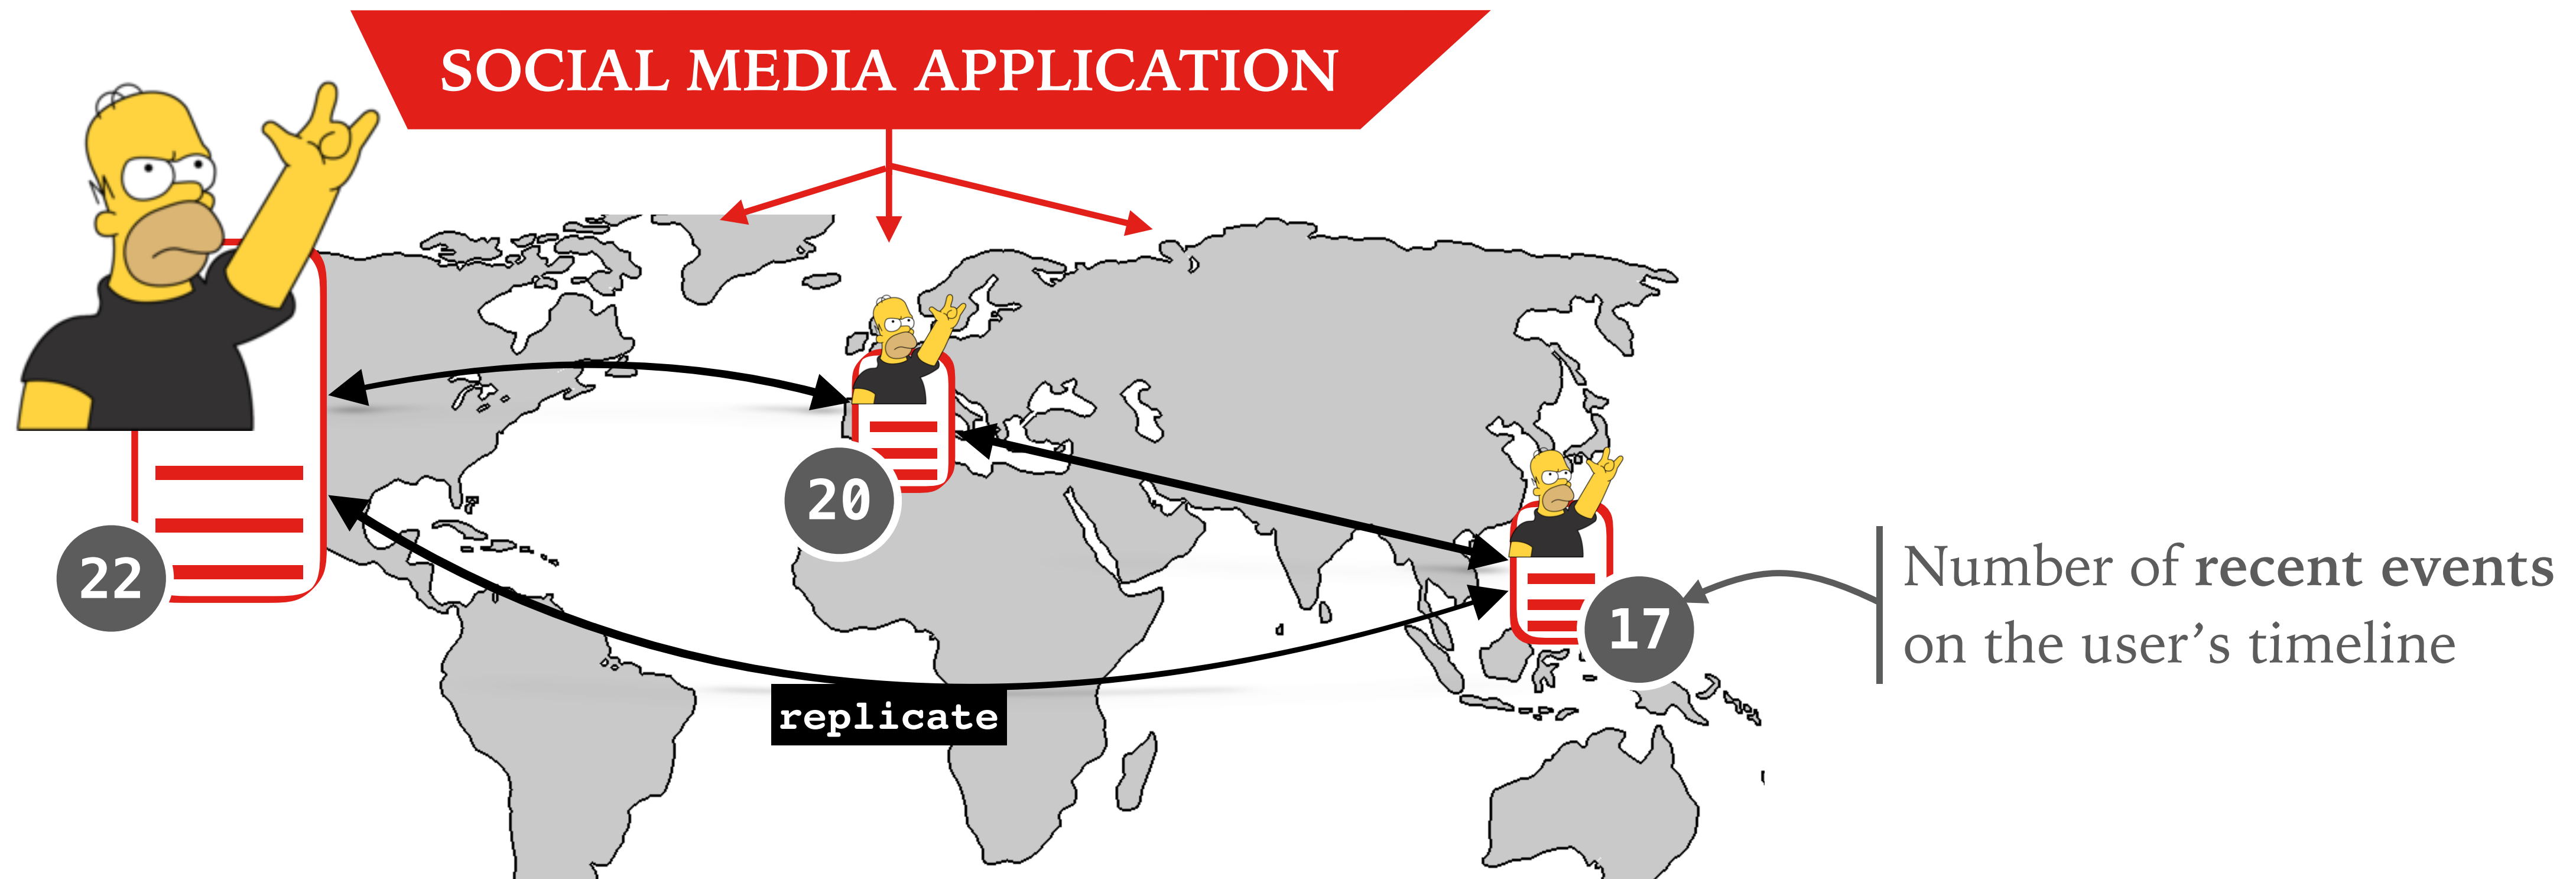






Strong Consistency

- Returns the **correct** data 22
 - Latency: ~200 ms
 - Can become **unavailable**
- [CAP], [PACELC]




Strong Consistency

- Returns the **correct** data 22
- Latency: ~200 ms
- Can become **unavailable**
[CAP], [PACELC]

Weak Consistency

- Latency: ~100 ms
- High availability
- Allows inconsistencies: can return


22 or 20 or 17



Strong Consistency

Weak Consistency

Neither model is ideal!



Strong Consistency

Weak Consistency


Neither model is ideal!



We use both models.

Multiple models

1. Weak consistency




2. Strong consistency



Multiple models

1. Weak consistency




2. Strong consistency



Increasingly many systems expose multiple consistency models:



App Engine



cassandra

Dynamo
[SOSP'07]




SimpleDB



Pileus
[SOSP'13]

Multiple models

1. Weak consistency



2. Strong consistency




Issues

1. *Send multiple requests?*
2. *How to leverage individual responses?*
3. *Semantics?*
4. *...*

Increasingly many systems expose multiple consistency models:



App Engine



cassandra

Dynamo
[SOSP'07]




SimpleDB



Pileus
[SOSP'13]

Multiple models

1. Weak consistency



2. Strong consistency




Issues

1. *Send multiple requests?*
2. *How to leverage individual responses?*
3. *Semantics?*
4. *...*

Increasingly many systems expose multiple consistency models:



App Engine



cassandra

Dynamo
[SOSP'07]



SimpleDB



Pileus
[SOSP'13]

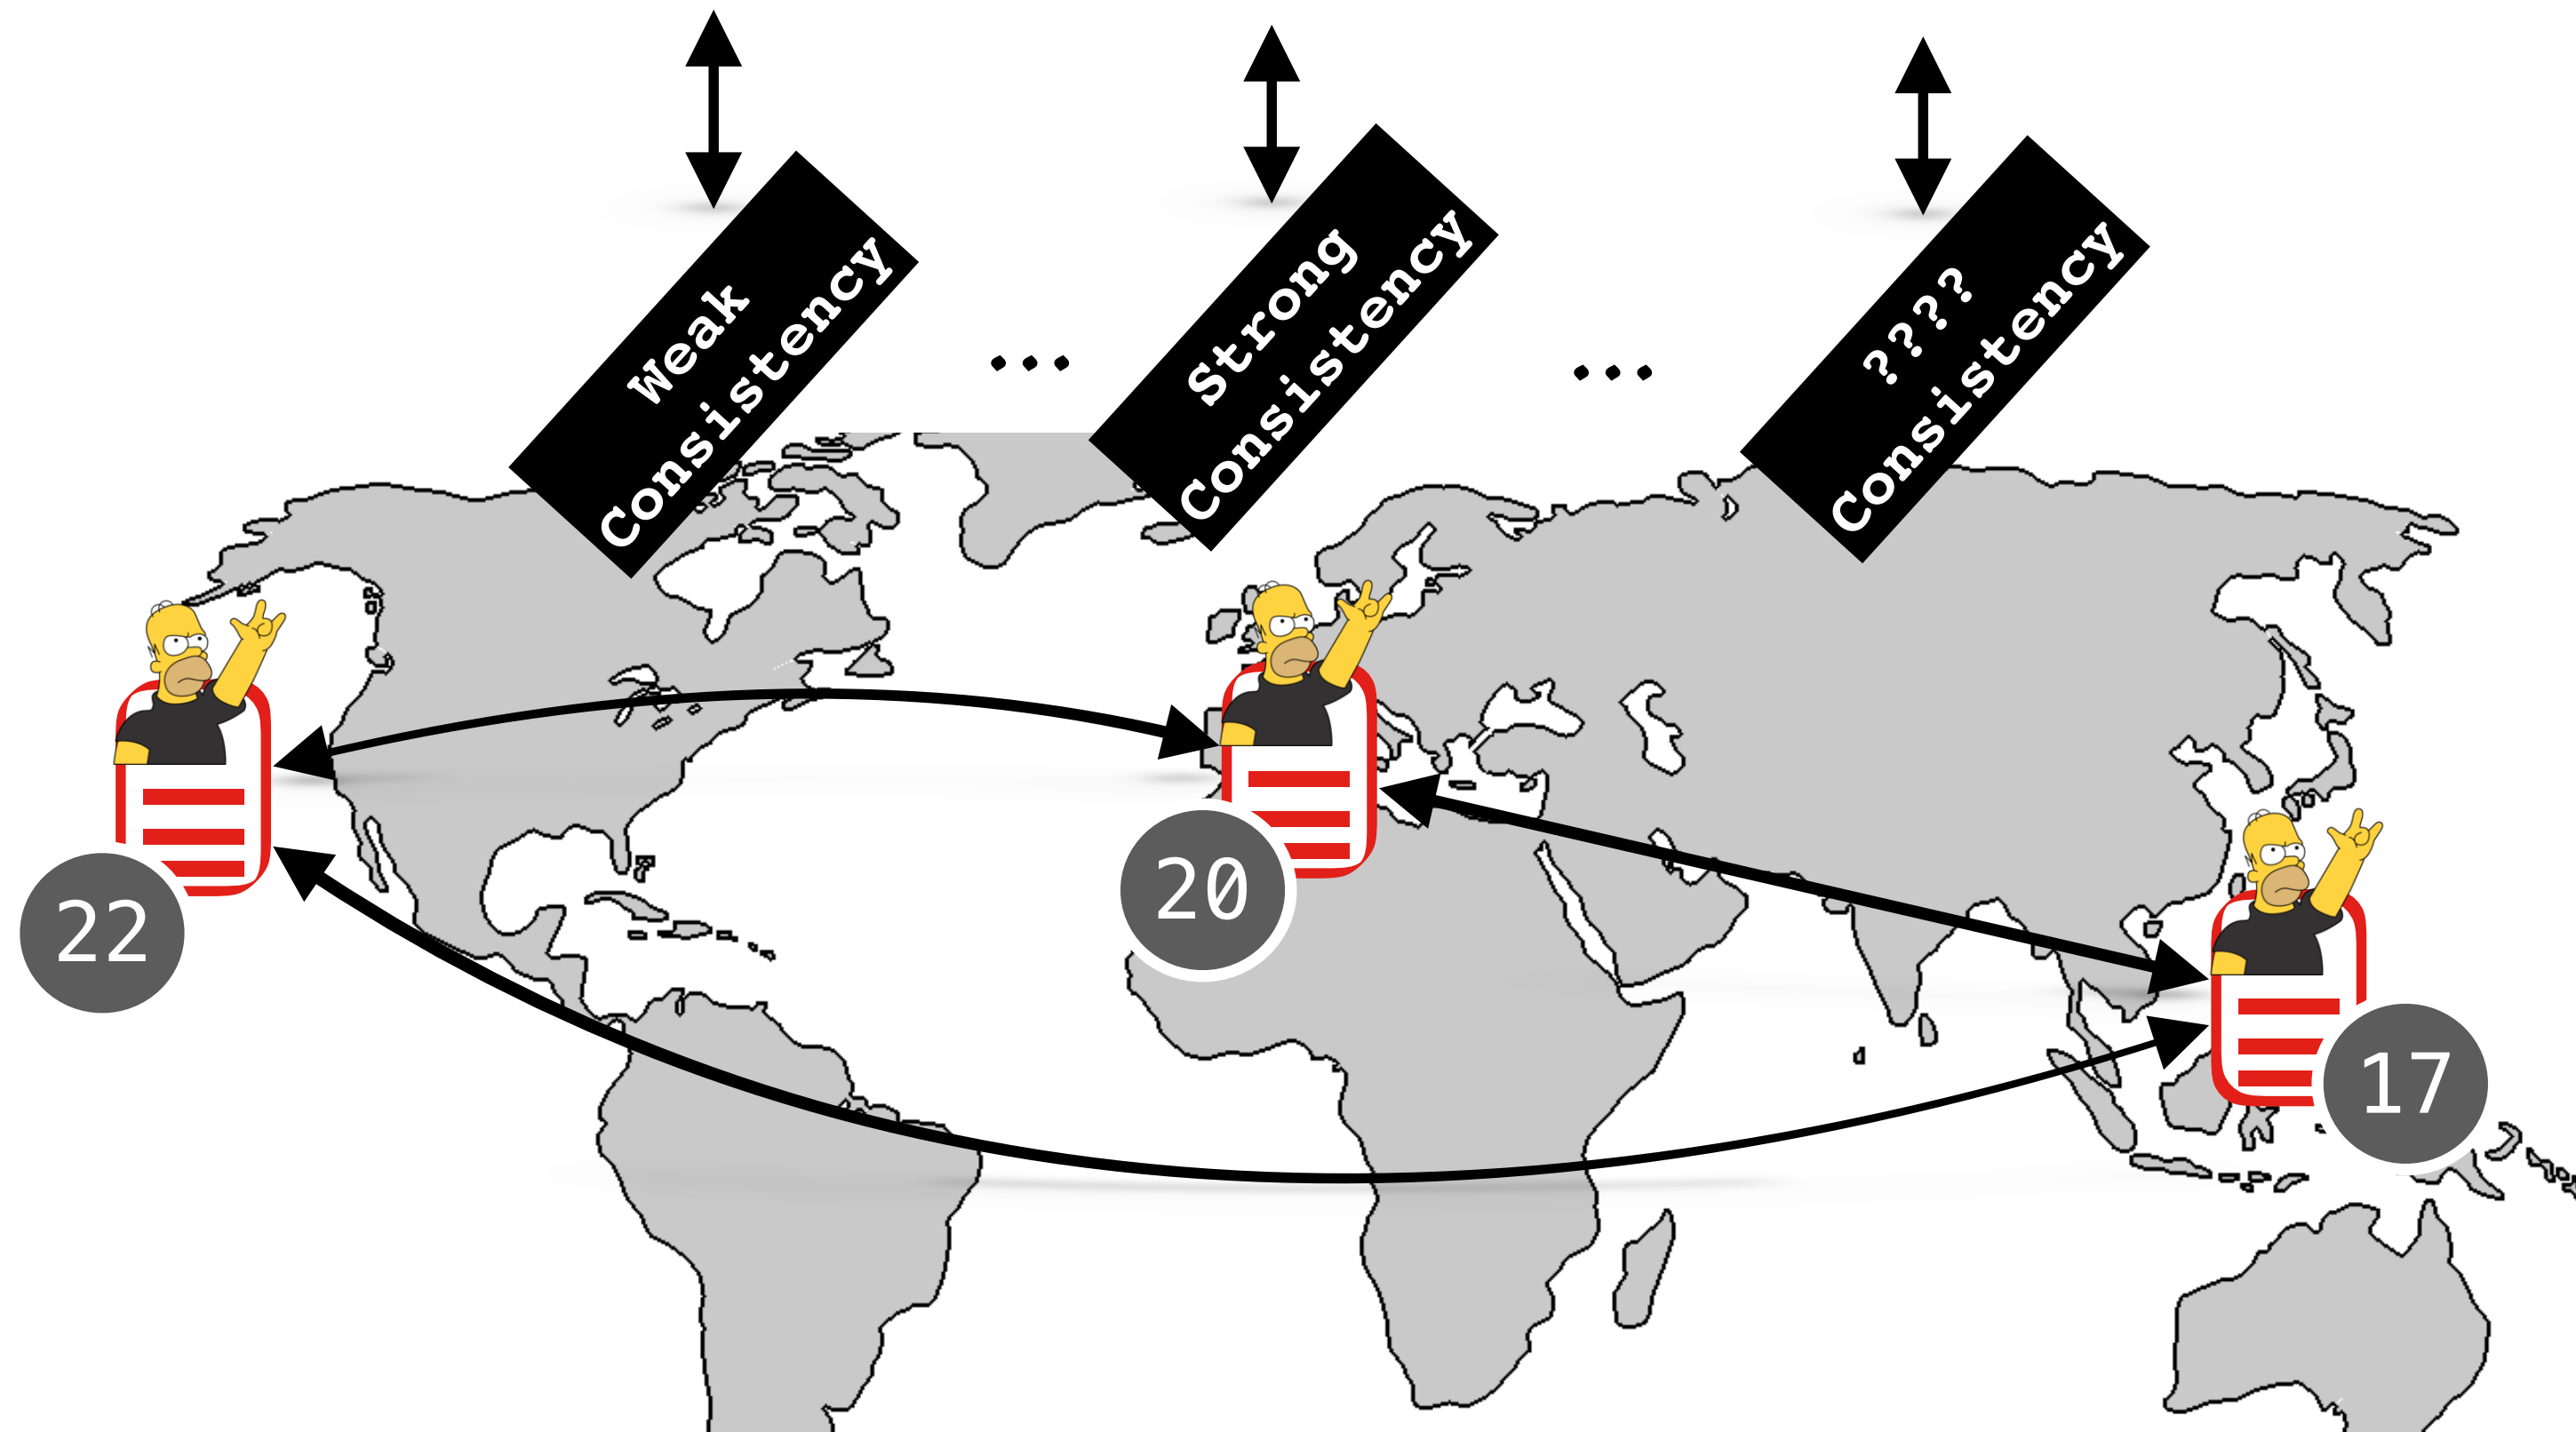
Problem

How do you program with

- ★ inconsistencies?
- ★ multiple values?

SOCIAL MEDIA APPLICATION

ABSTRACTION
FOR REPLICATED OBJECTS



SOCIAL MEDIA APPLICATION


**ABSTRACTION
FOR REPLICATED OBJECTS**

**Weak
Consistency**

**Strong
Consistency**

**???
Consistency**


CORRECTABLE




SOCIAL MEDIA APPLICATION




**ABSTRACTION
FOR REPLICATED OBJECTS**



**Weak
Consistency**



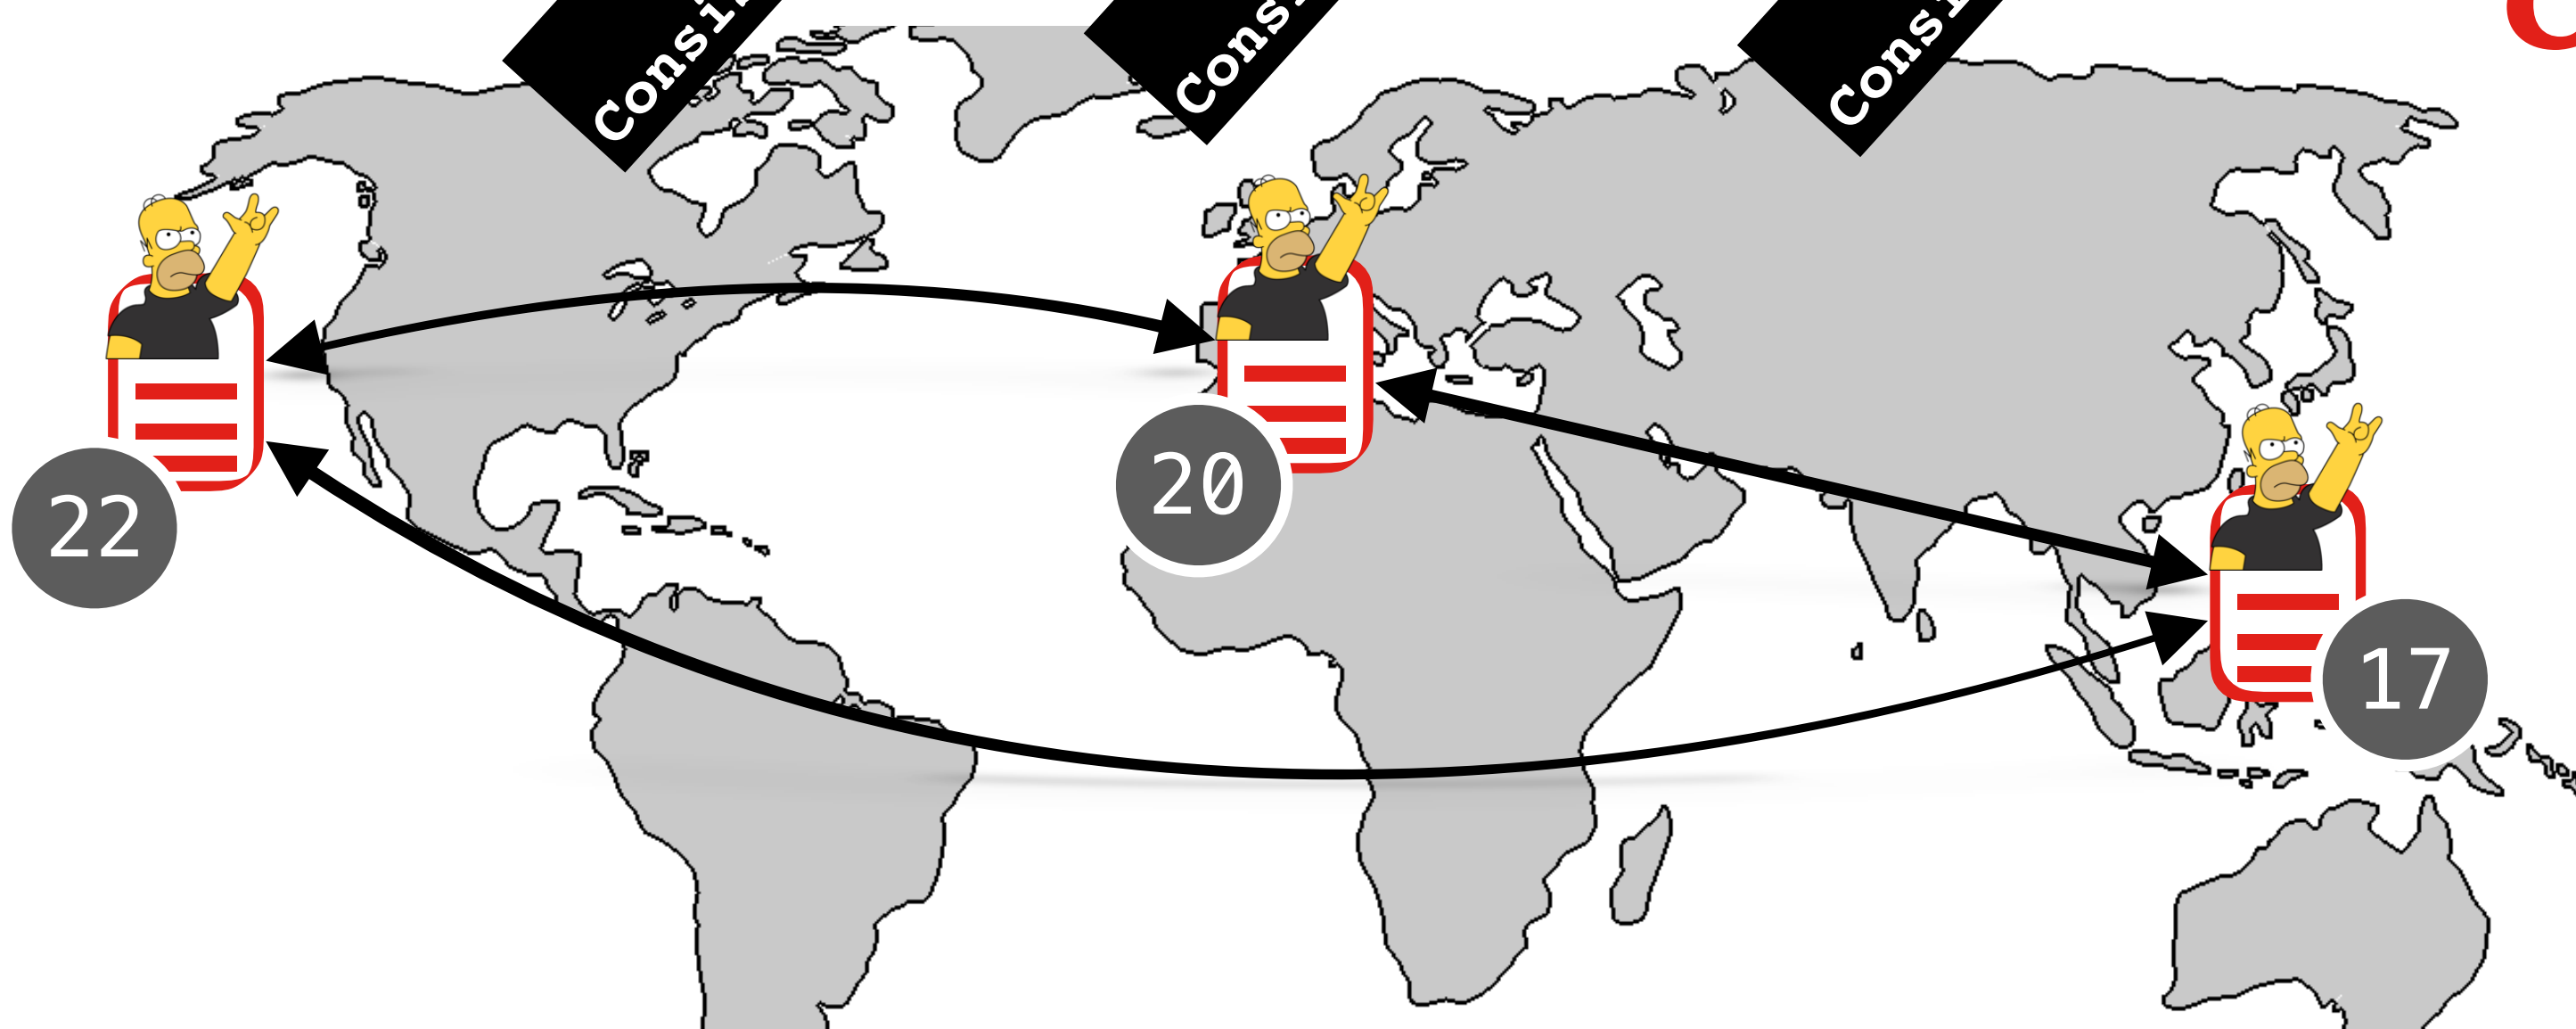
**Strong
Consistency**




**???
Consistency**

...

...



**Incremental
Consistency
Guarantees (ICG)**




provides

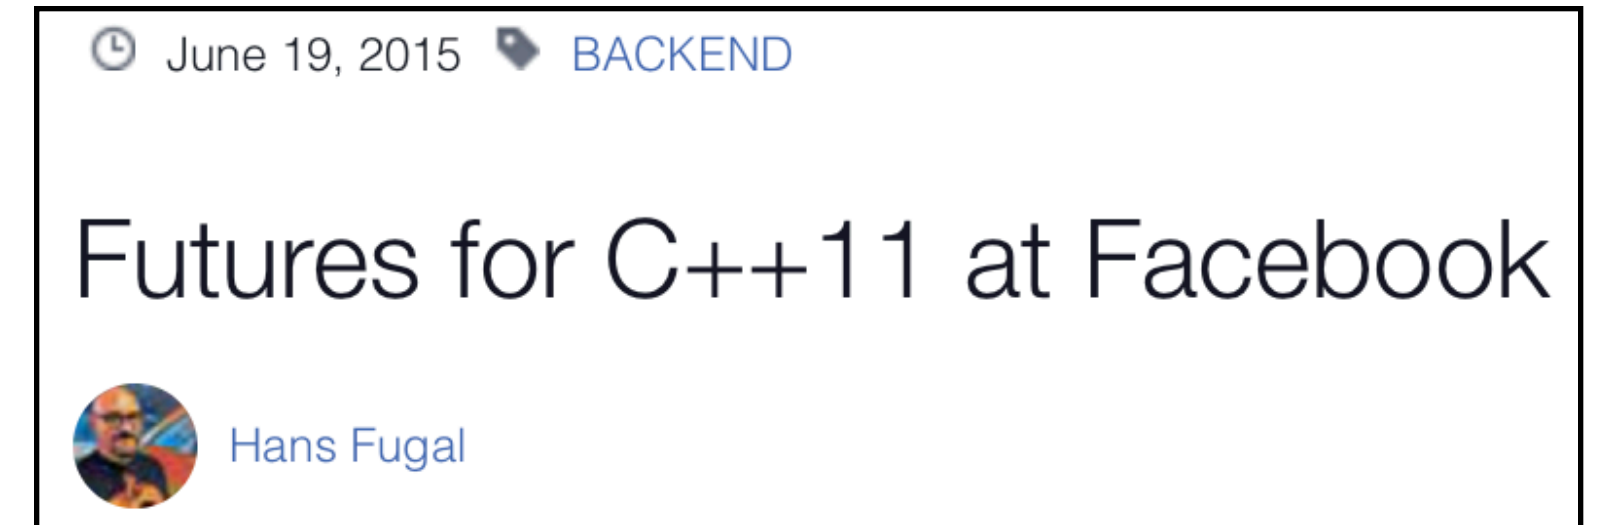

CORRECTABLE

Correctables / Design

- Starting point: **Promises**
 - Placeholders for values
 - Becoming mainstream




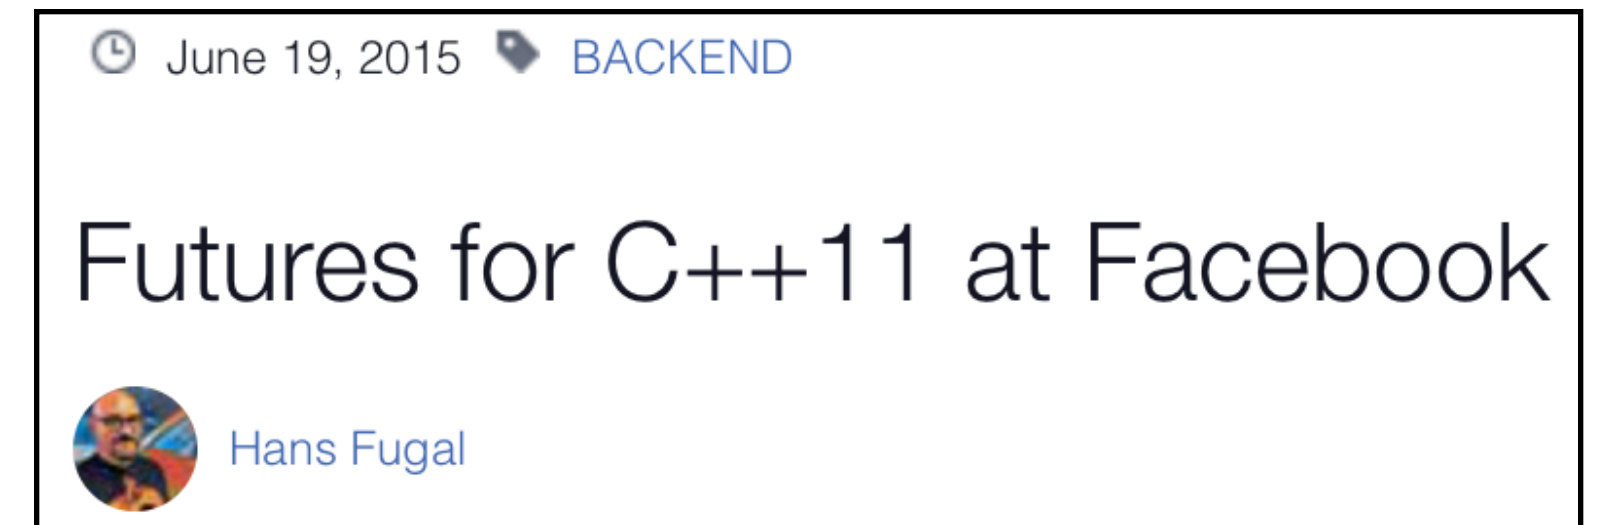
Futures and Promises



Correctables / Design

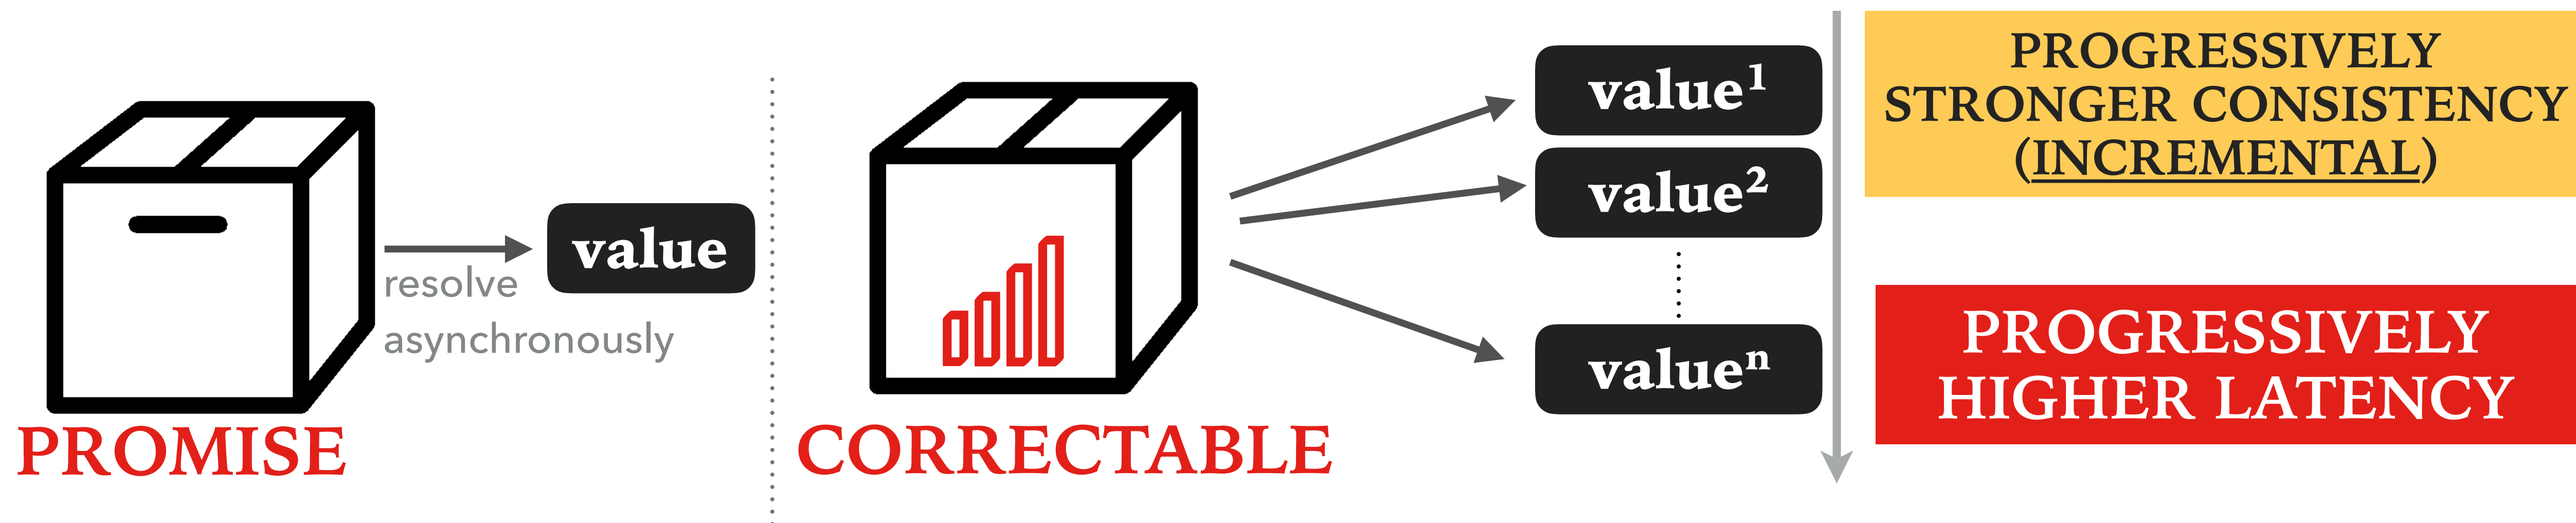

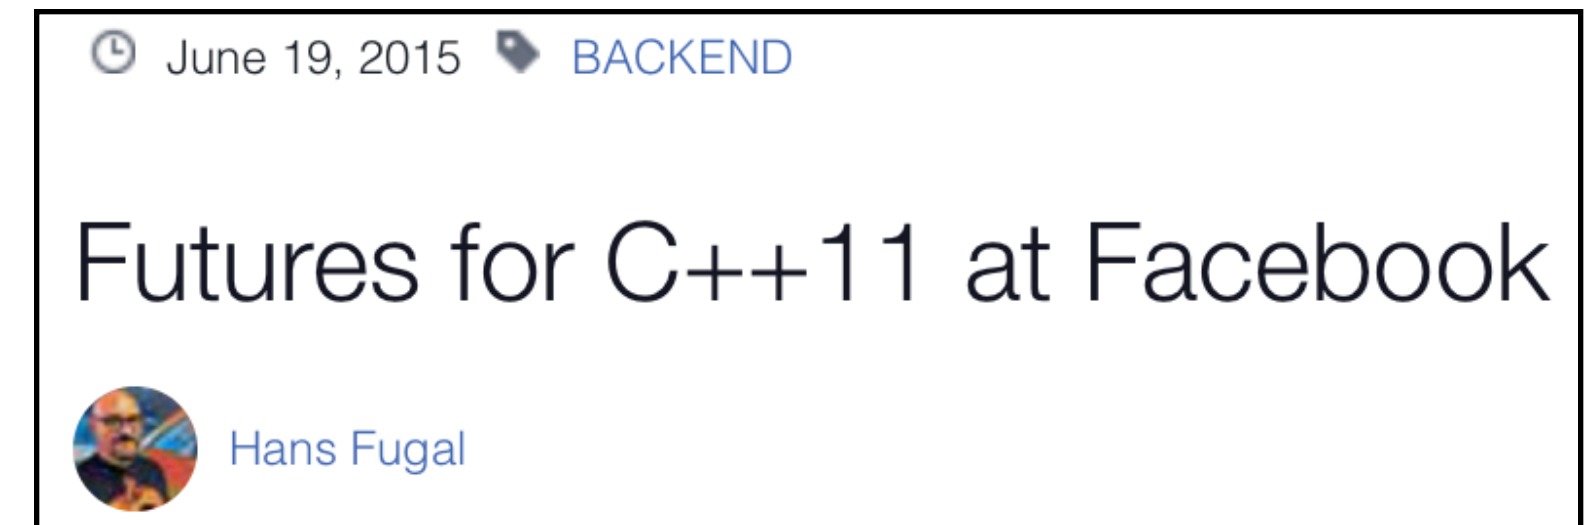

- Starting point: **Promises**
 - Placeholders for values
 - Becoming mainstream

Futures and Promises




Correctables / Design

Futures and Promises



Consistency Models are **Complementary**



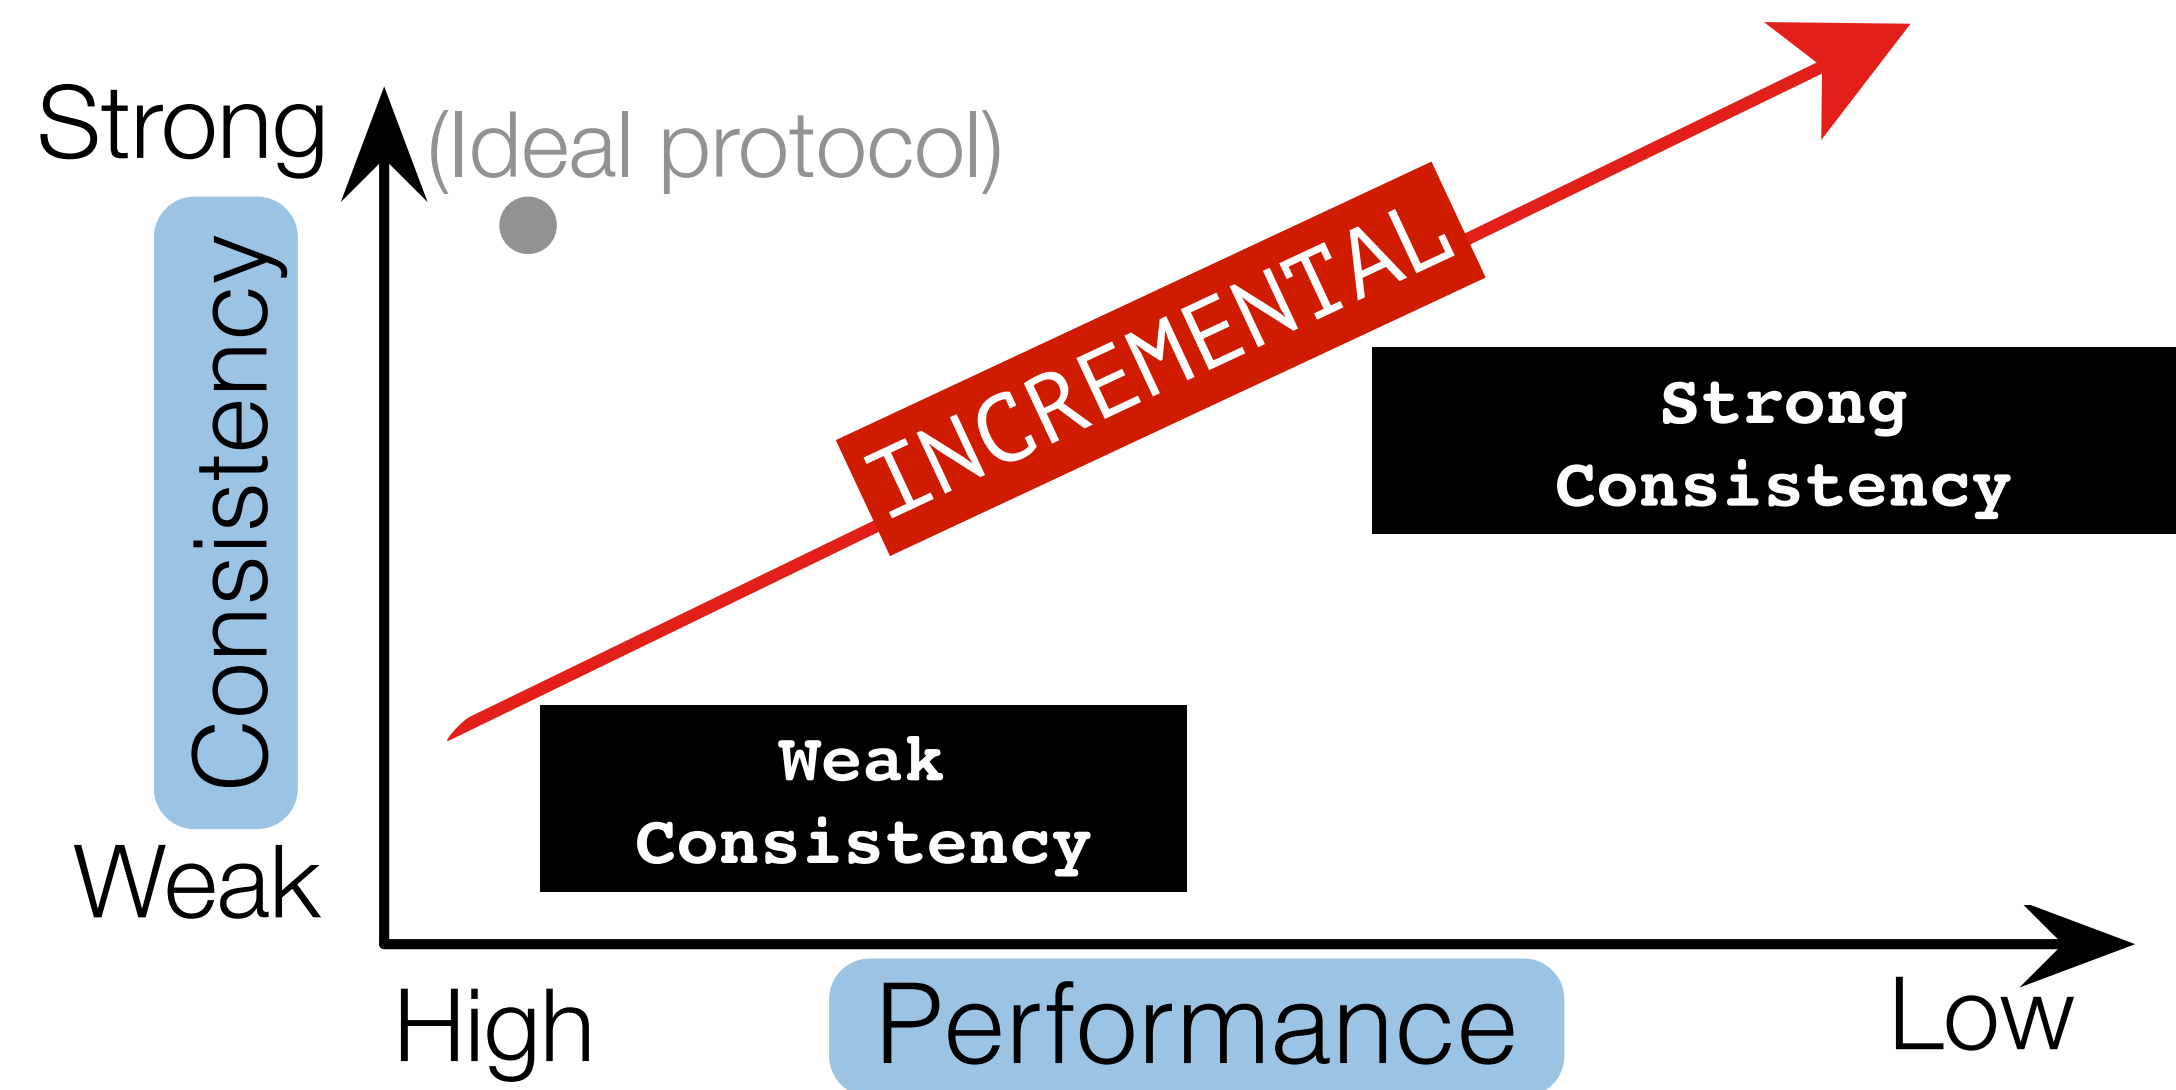
Weak consistency:

- ★ Fast
- ★ (Often correct)

Strong consistency:

- ★ Slower
- ★ (Correct with certainty)

Consistency Models are **Complementary**



Weak consistency:


- ★ Fast
- ★ (Often correct)

Strong consistency:

- ★ Slower
- ★ (Correct with certainty)

So what?

Consistency Models are **Complementary**



Weak consistency:

- ★ Fast
- ★ (Often correct)

Strong consistency:

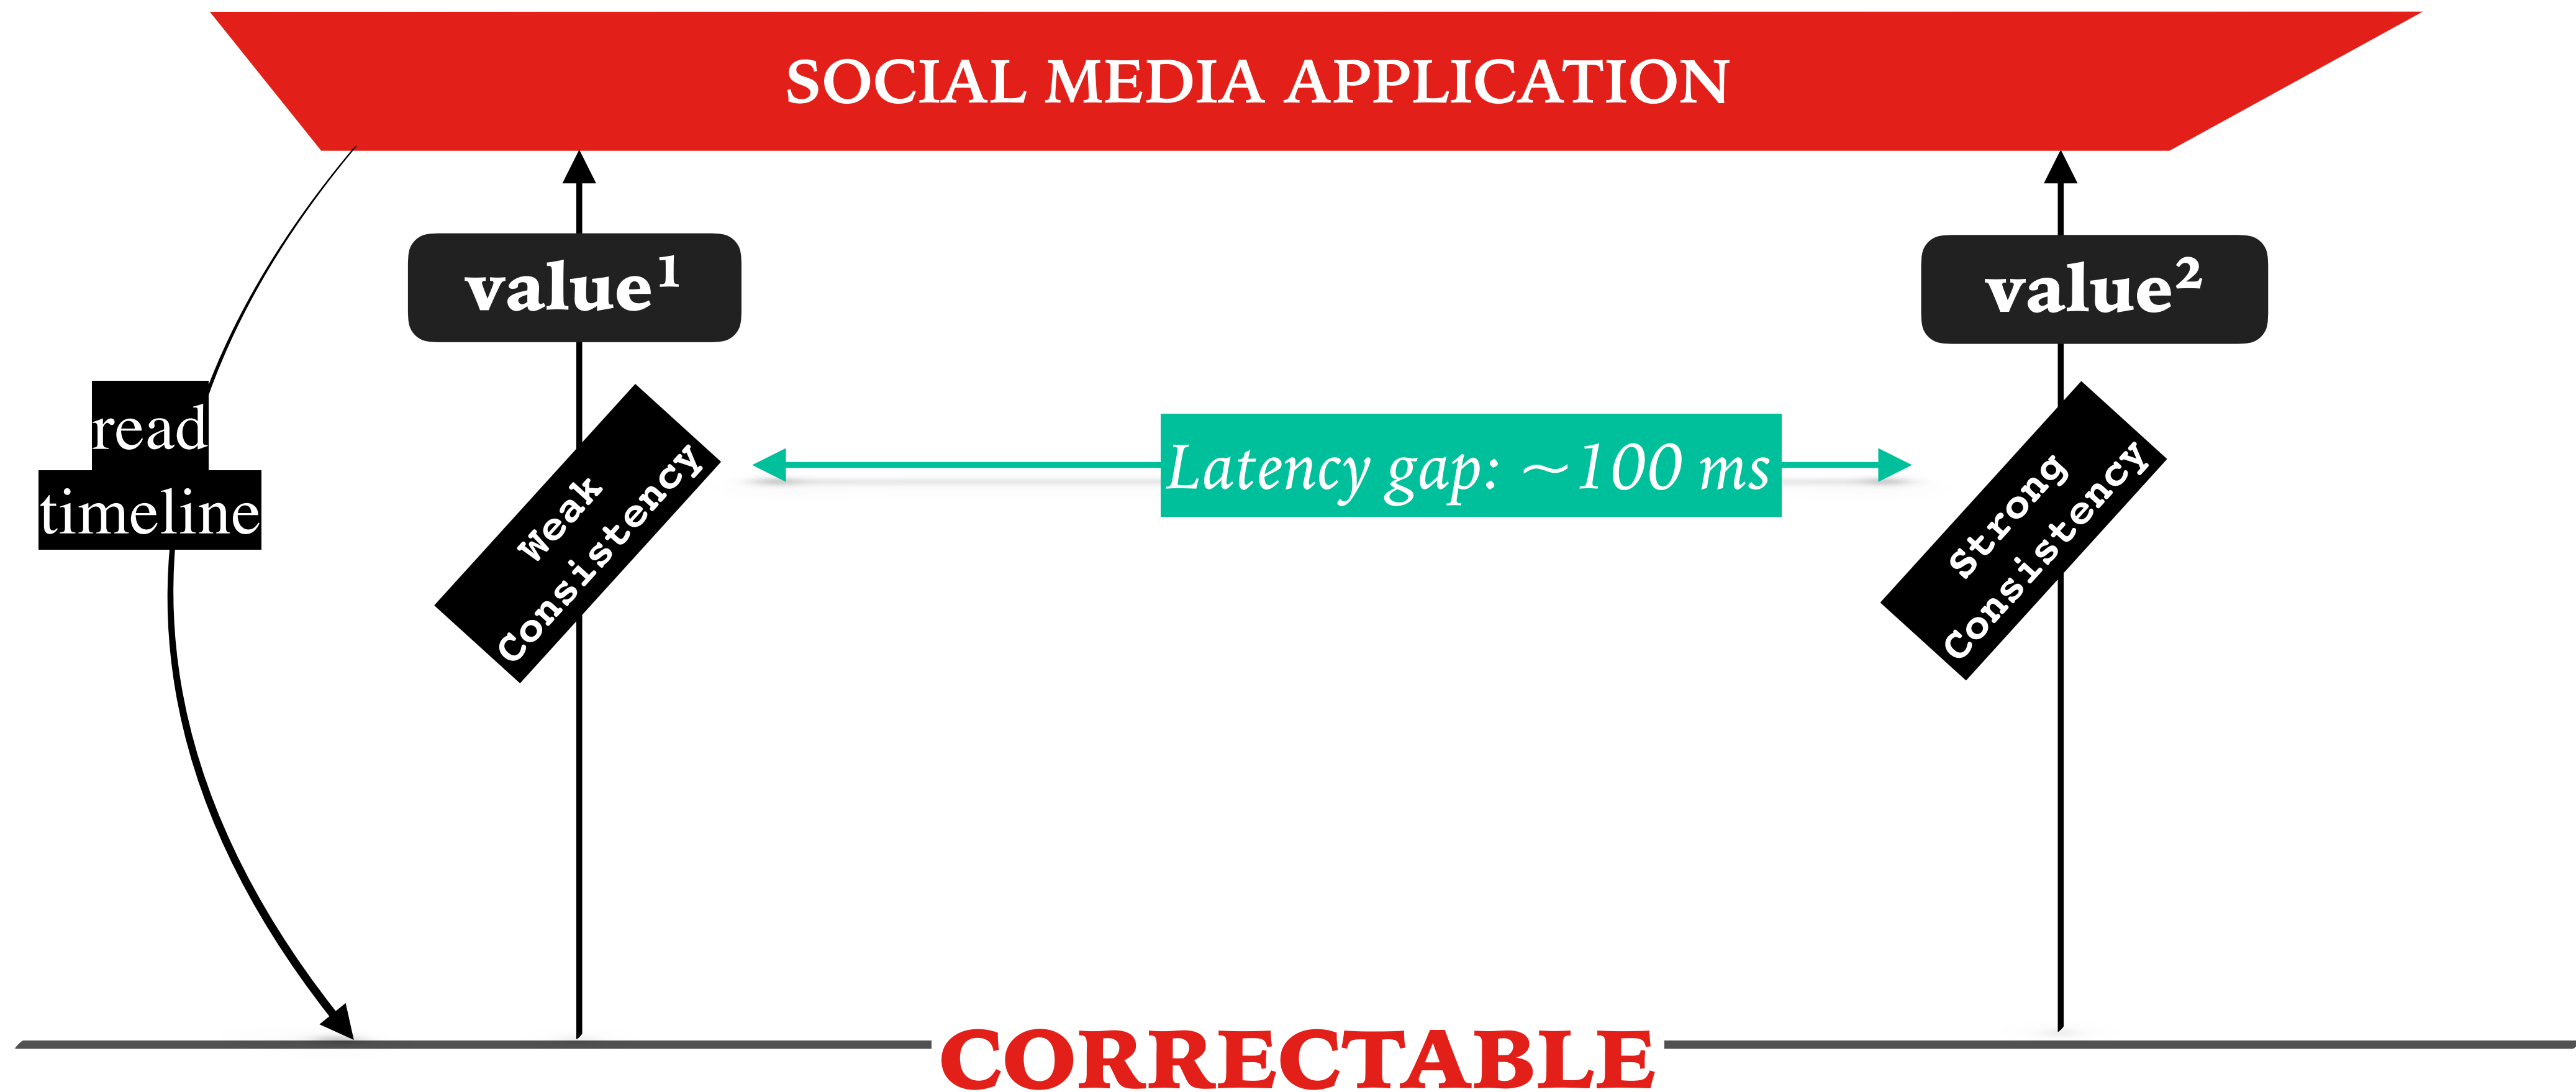
- ★ Slower
- ★ (Correct with certainty)

So what?
Latency optimizations

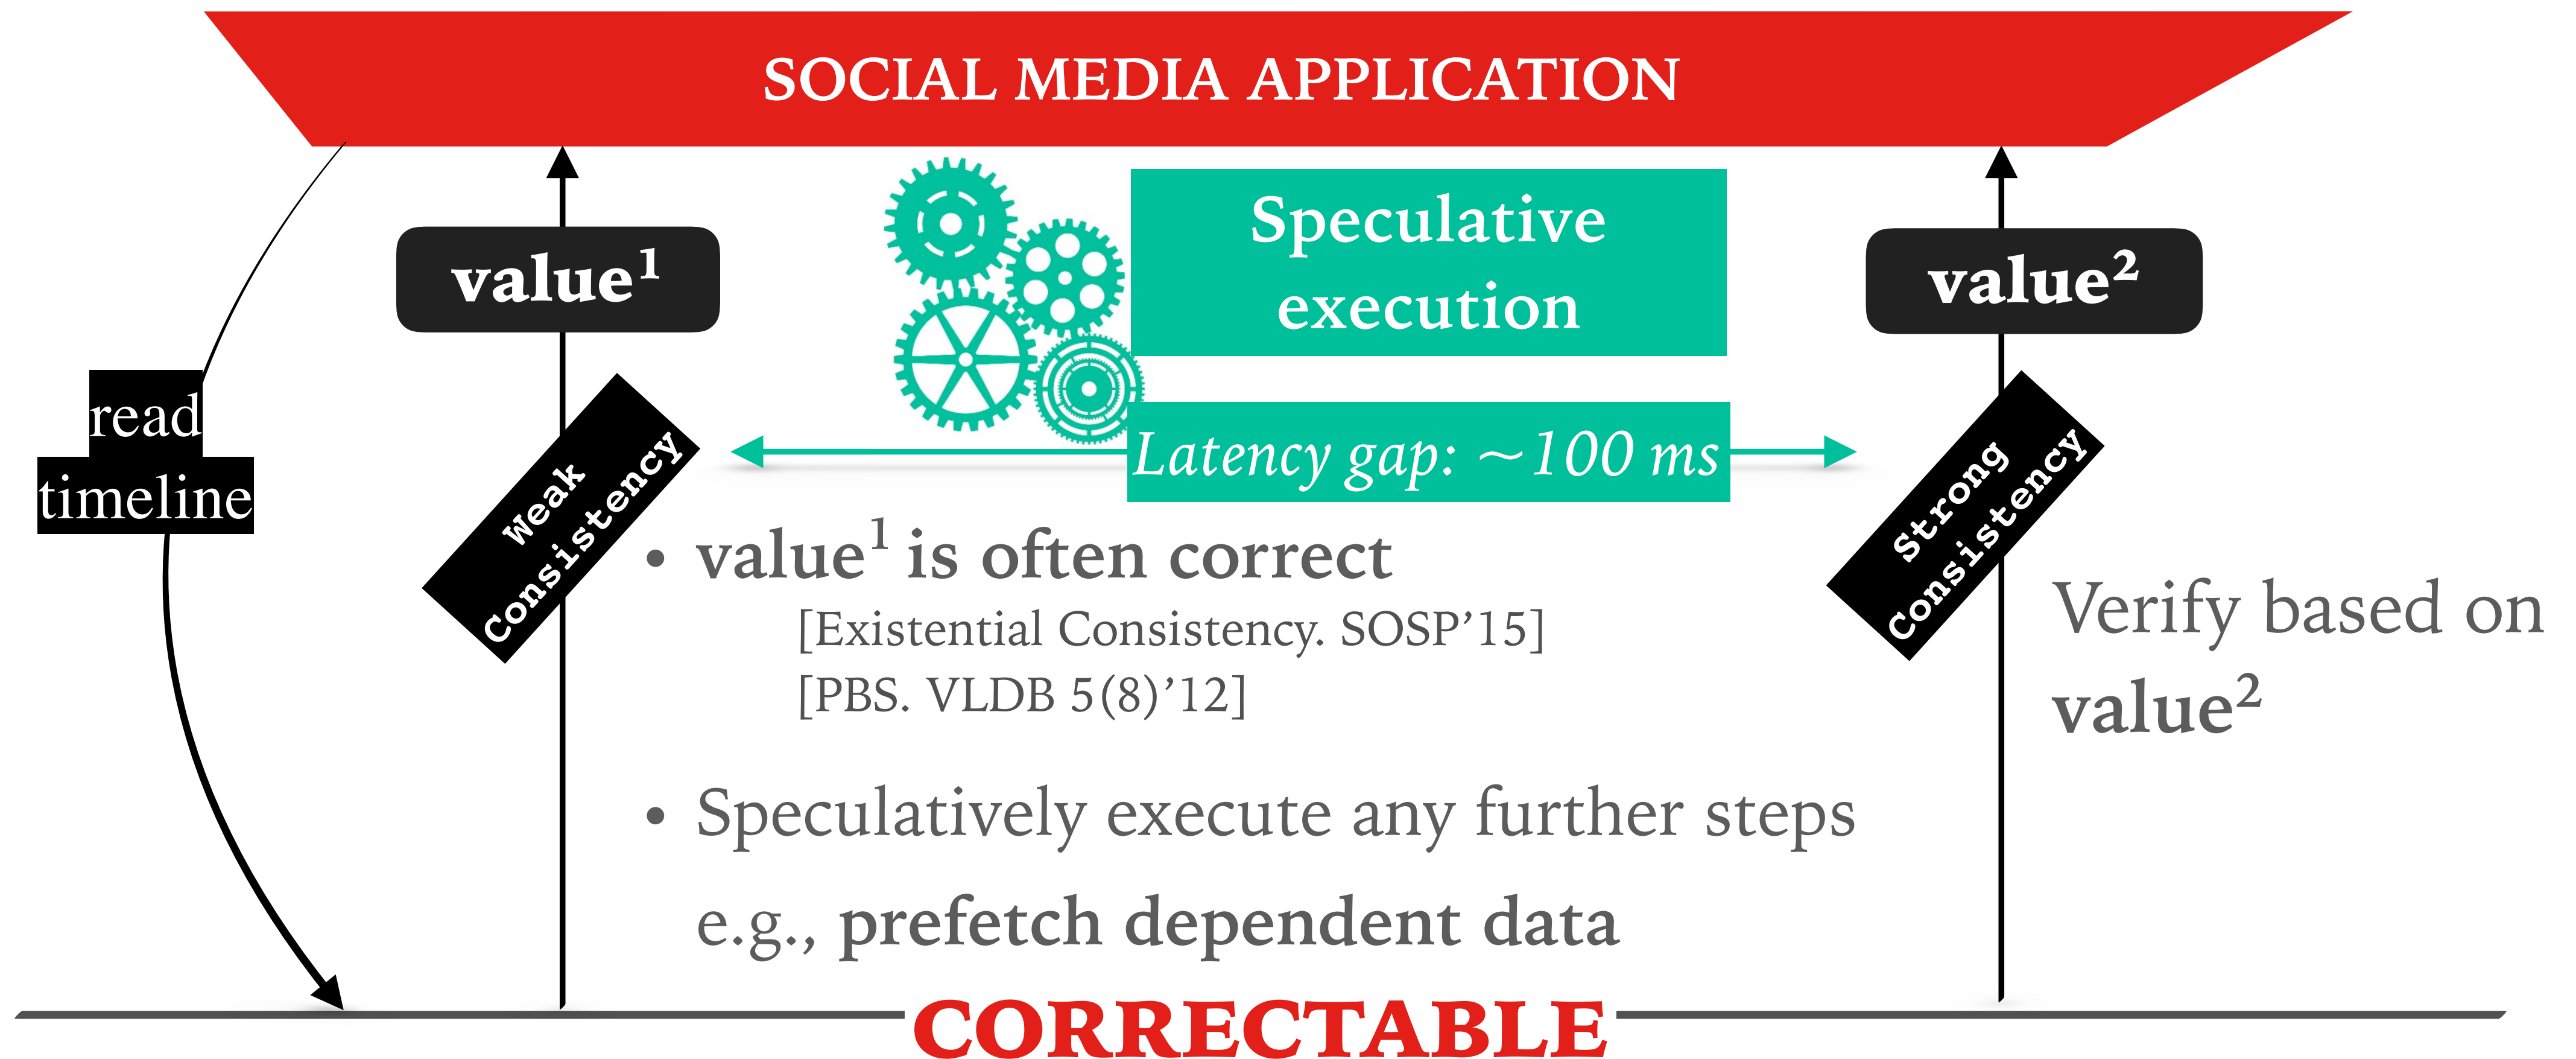
Speculating with Correctables



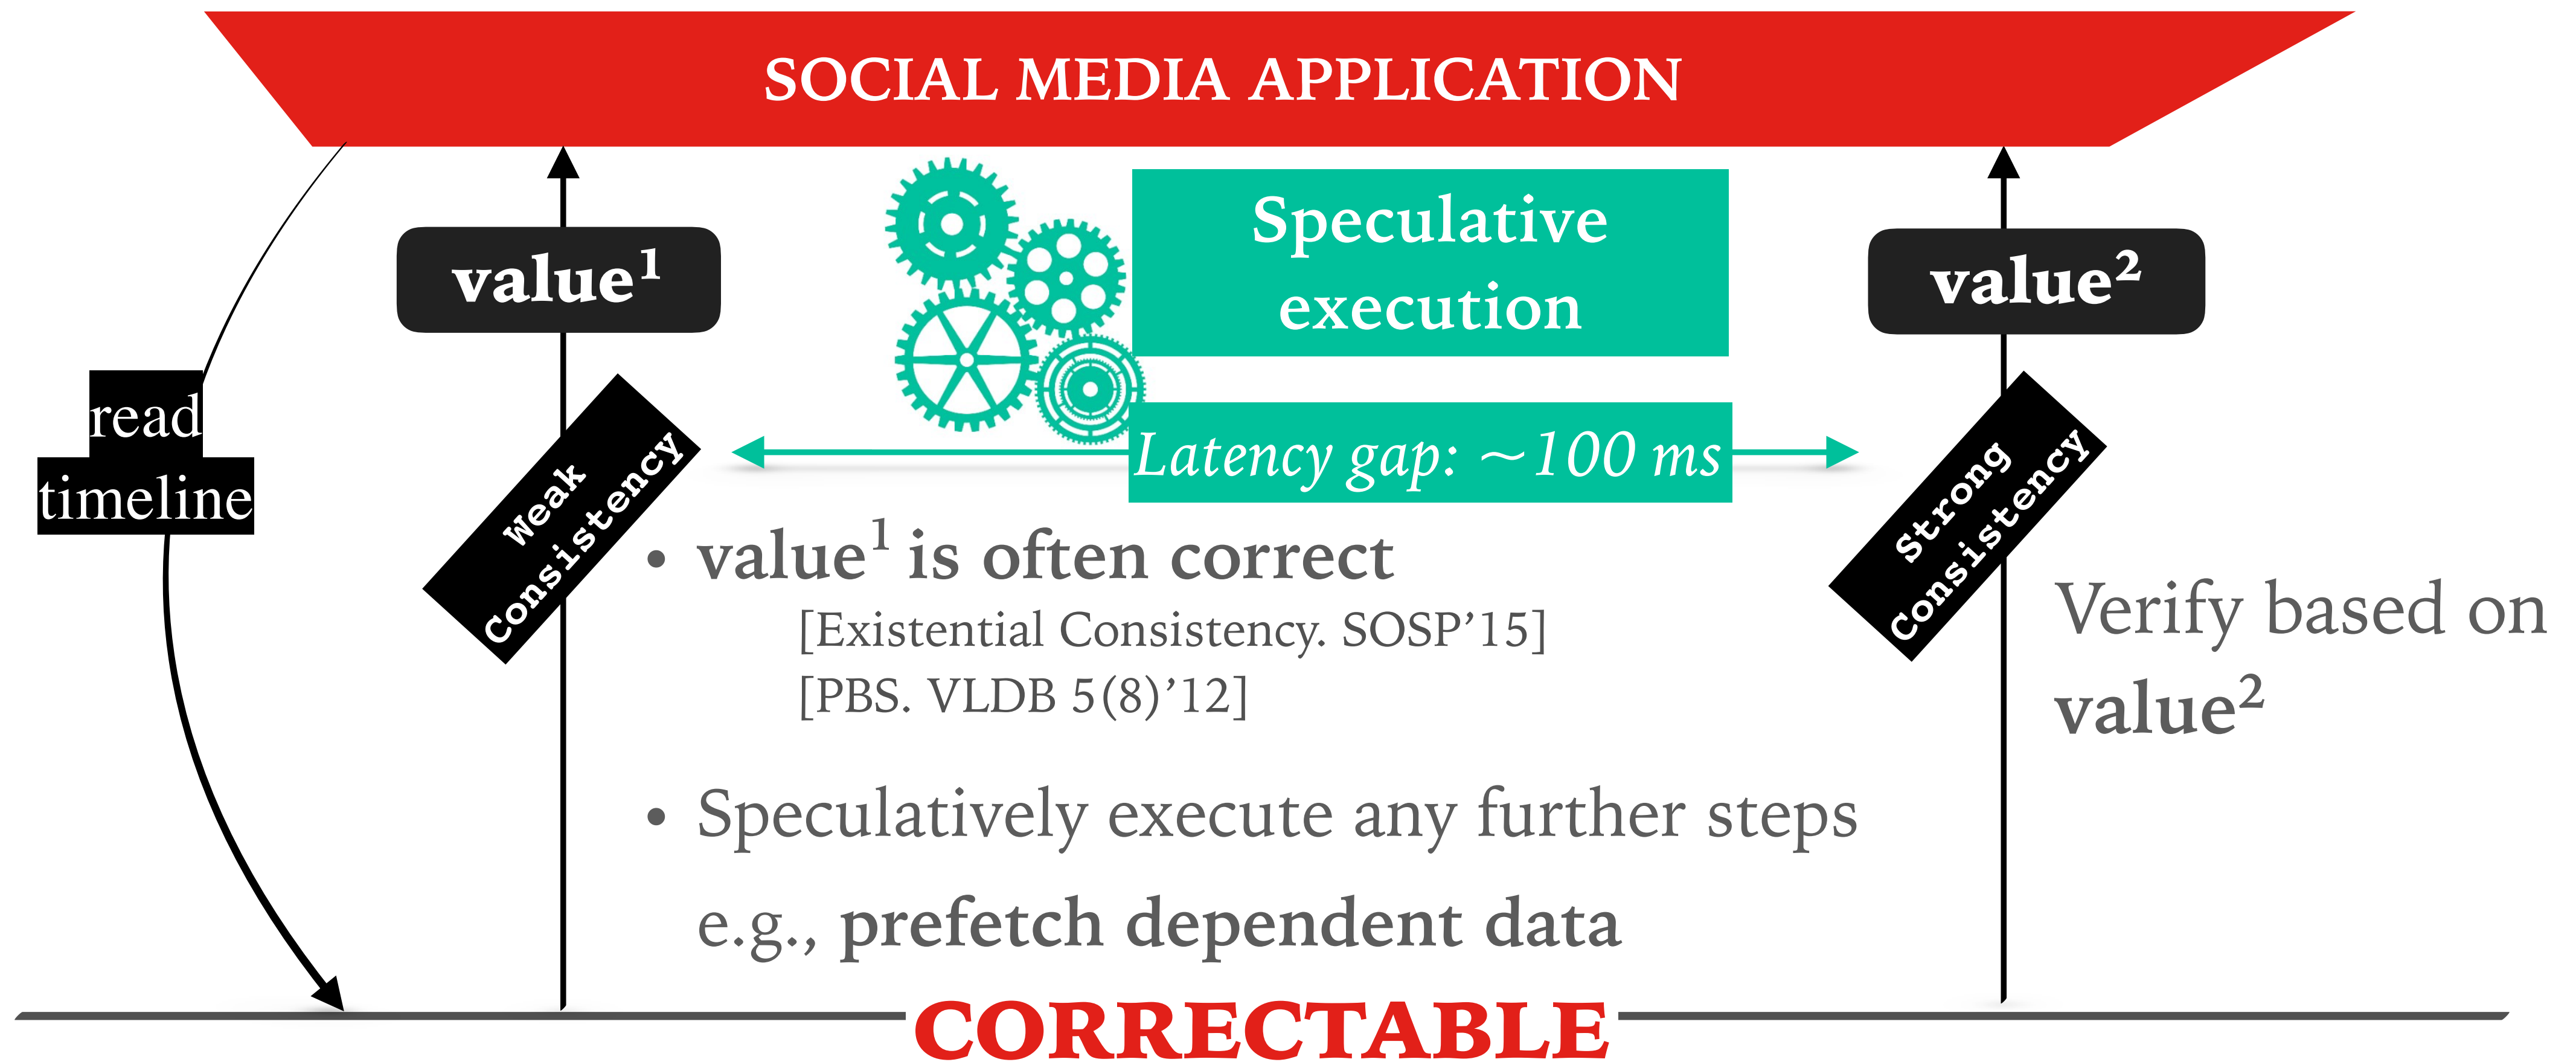
Speculating with Correctables



Speculating with Correctables




Speculating with Correctables




Lower latency of strong consistency

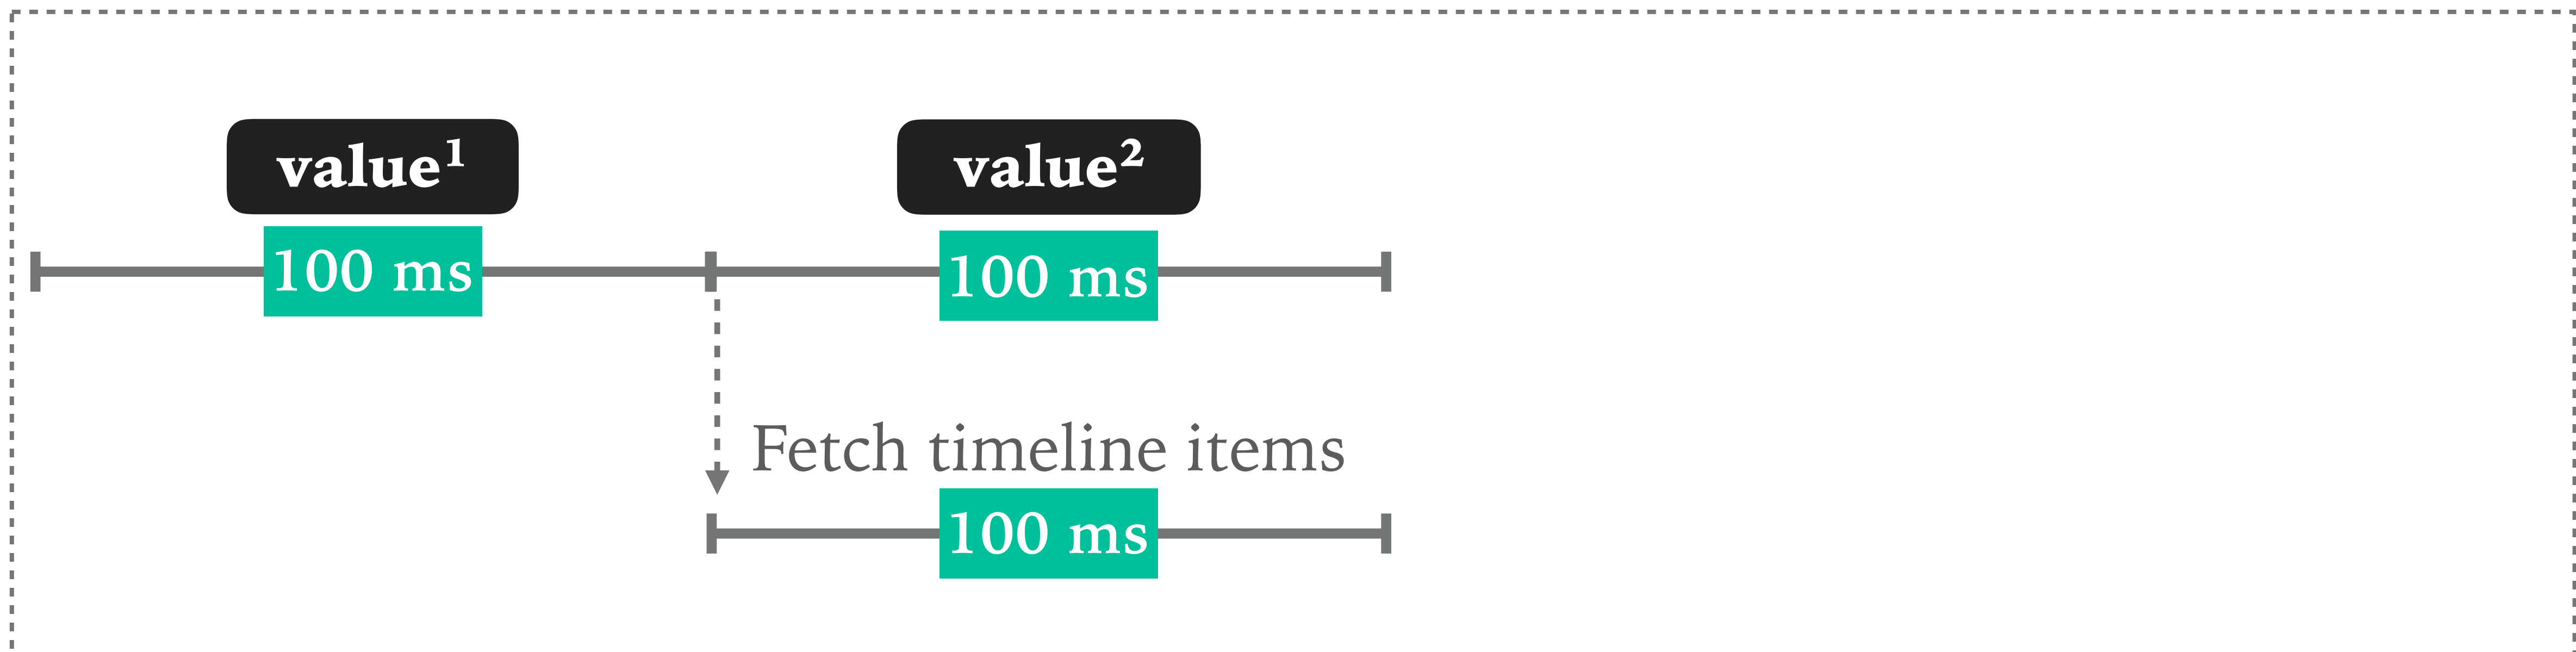
Traditional operation:



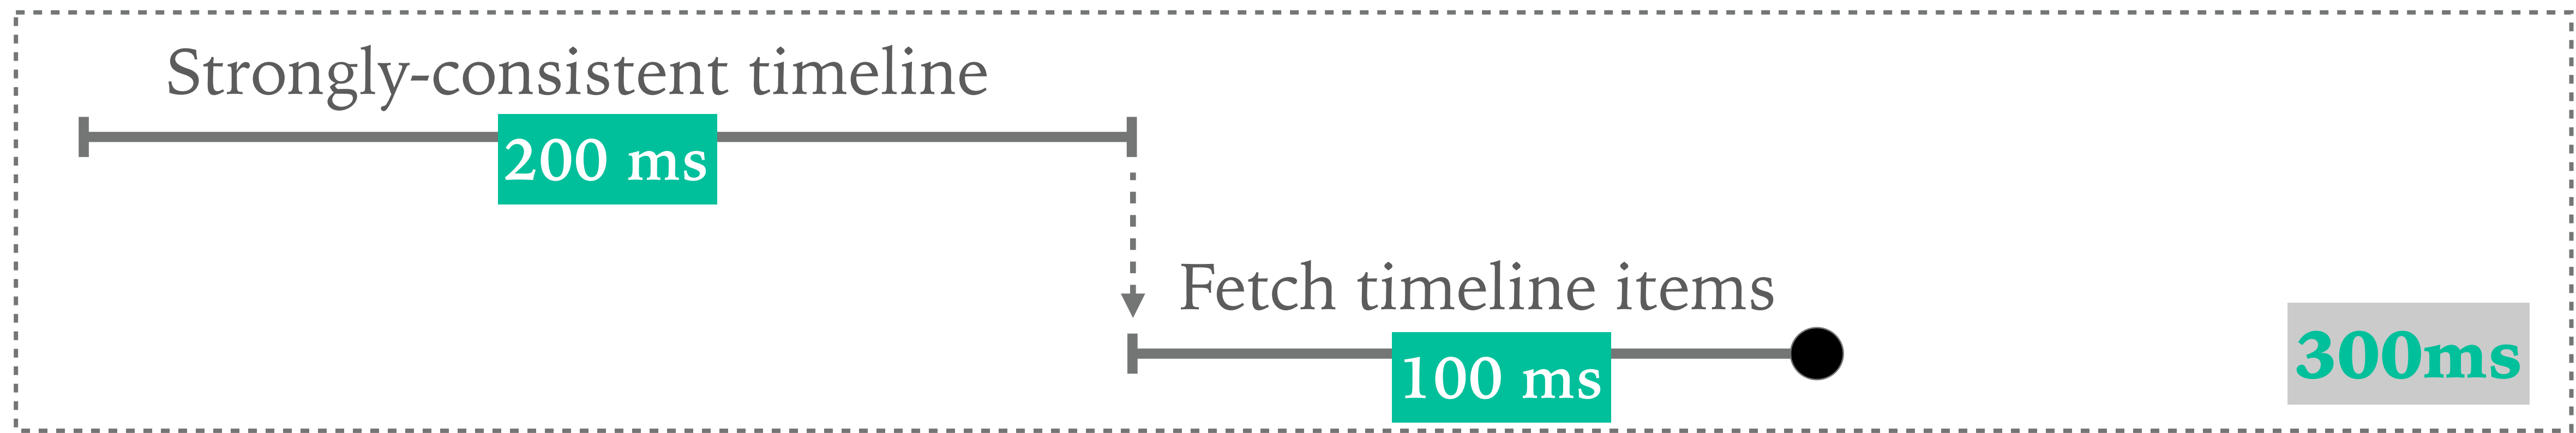
Traditional operation:



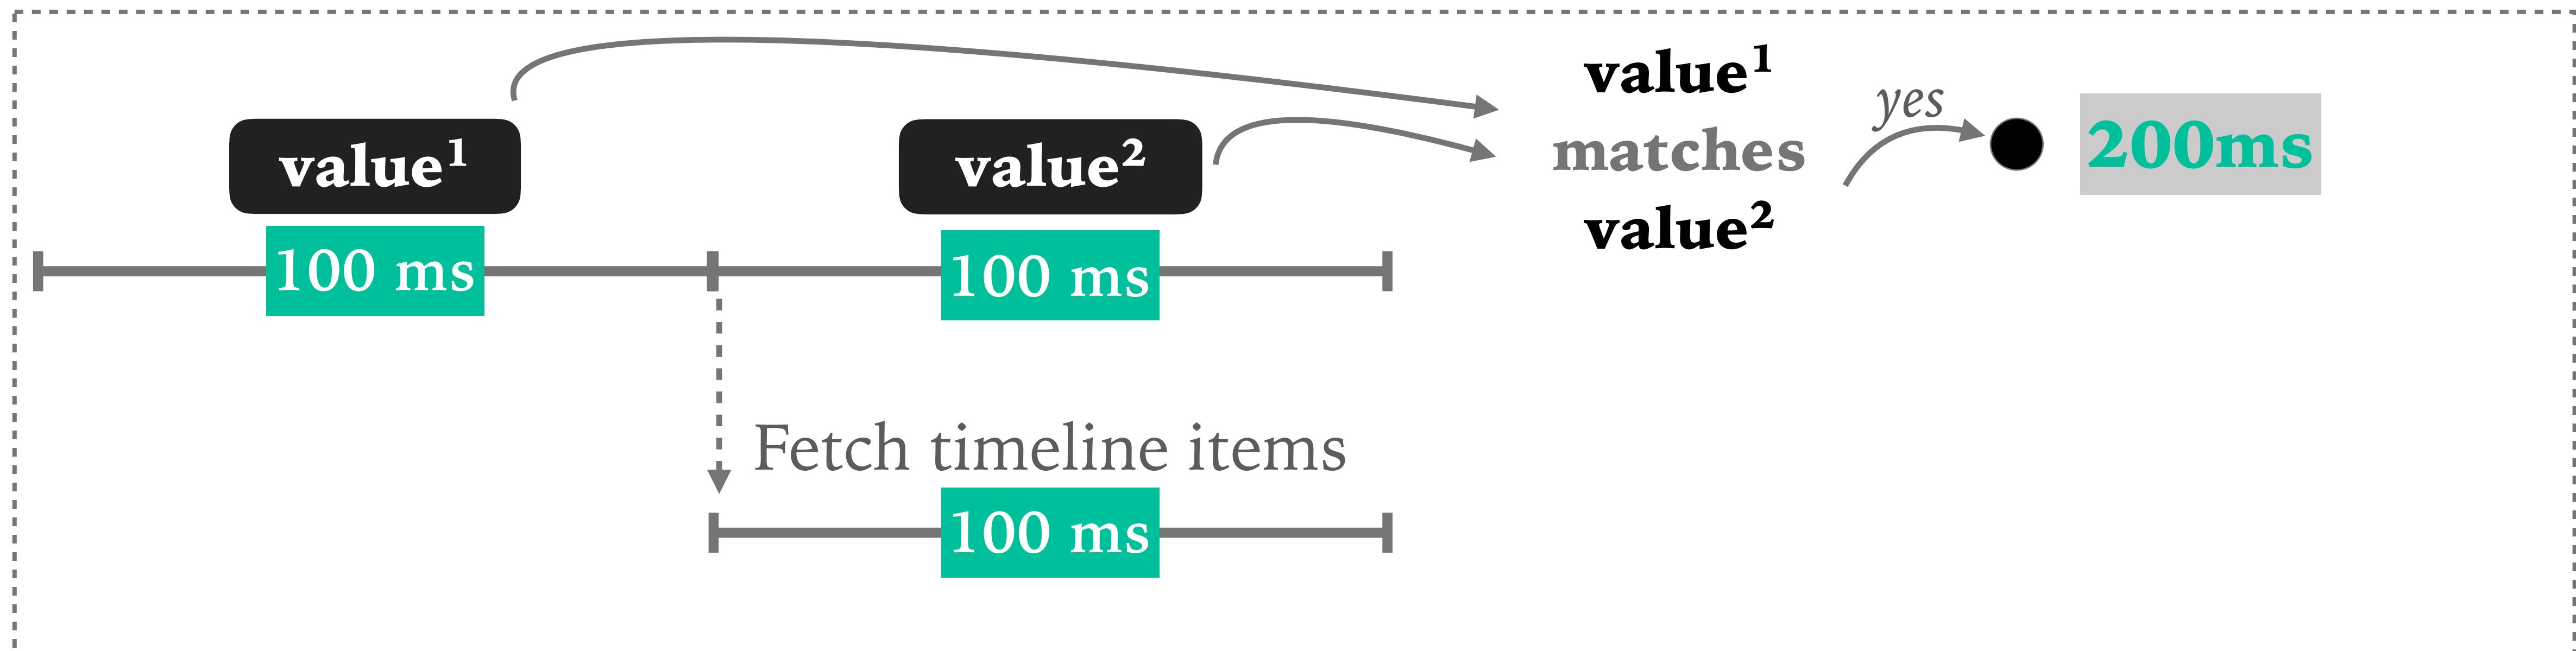
Speculative operation with ICG:




Traditional operation:



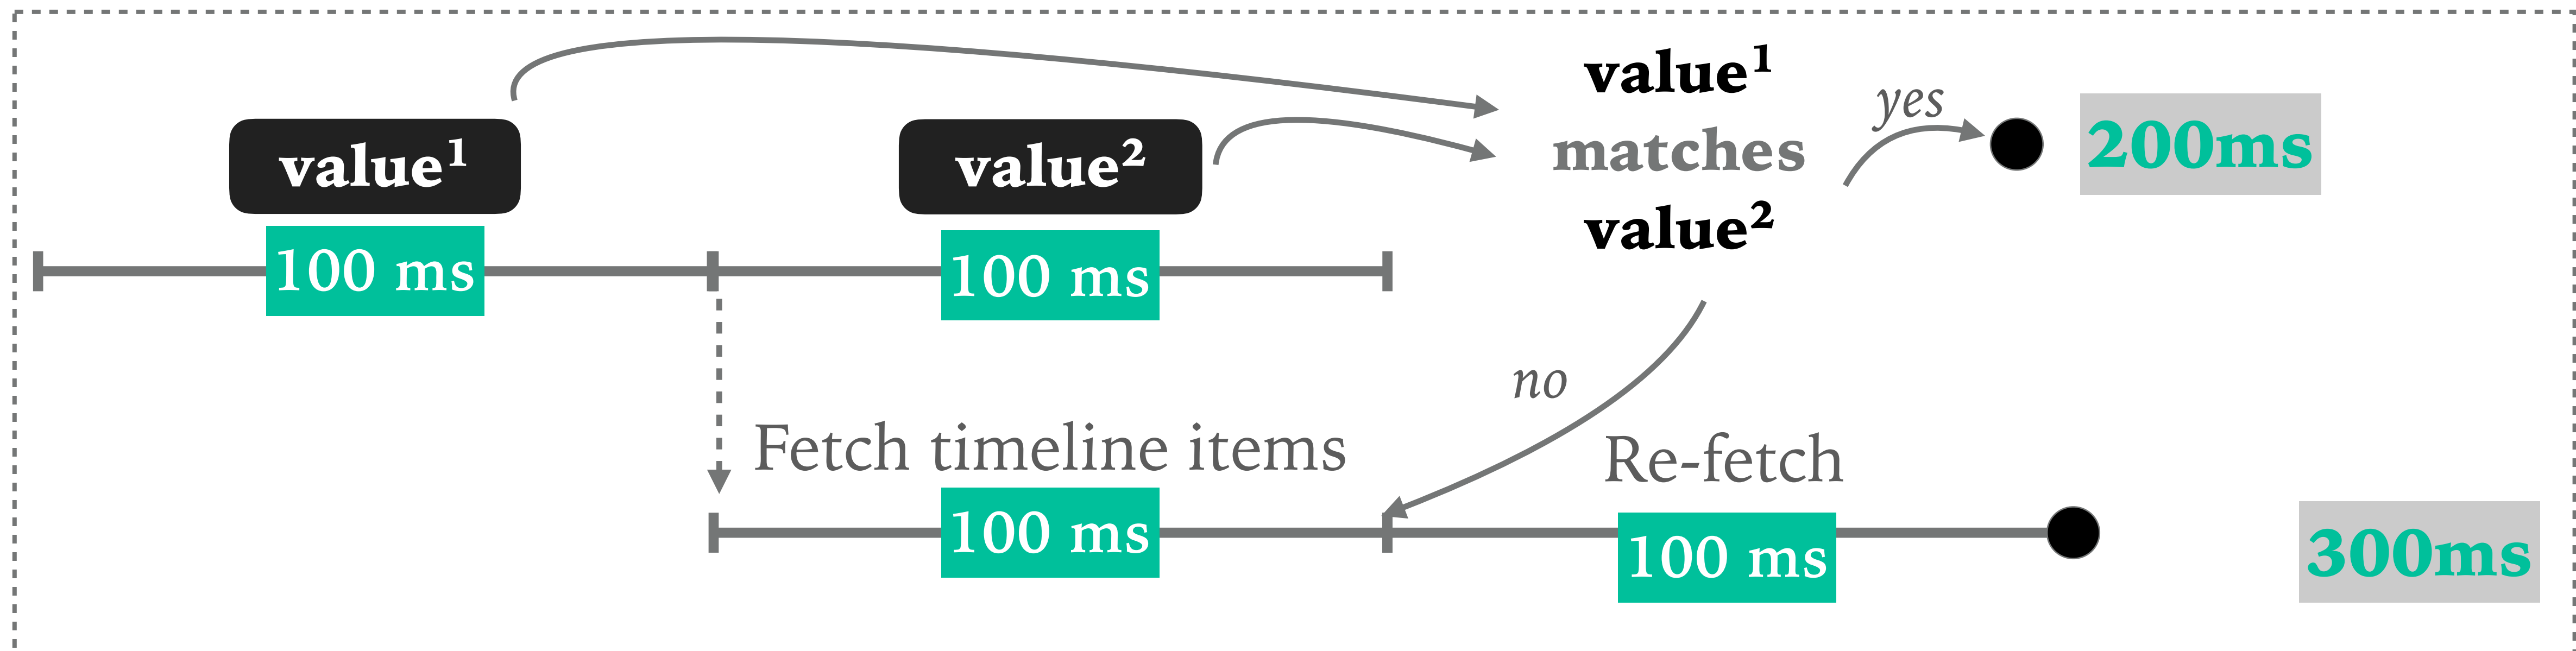
Speculative operation with ICG:



Traditional operation:



Speculative operation with ICG:



Speculation case-study

- Application: **Twissandra**
- Workload generated via YCSB
- Clients in Ireland
- Geo-replication on Amazon's EC2



Speculation case-study

- Application: **Twissandra**
- Workload generated via YCSB
- Clients in Ireland
- Geo-replication on Amazon's EC2



Incremental Consistency Guarantees for Replicated Objects

Rachid Guerraoui, Matej Pavlovic, and Dragos-Adrian Seredinschi*

*School of Computer and Communication Sciences,
École Polytechnique Fédérale de Lausanne (EPFL), Switzerland
{rachid.guerraoui, matej.pavlovic, dragos-adrian.seredinschi}@epfl.ch*

check the paper

- ★ Advertising System
— Speculation case-study
- ★ Ticket-selling System
— Exploiting application semantics
- ★ Overheads evaluation
& Optimizations
- ★ Latency gaps between consistency models

Decreasing latency of strong consistency

What is the
latency of the **fetch_timeline()** operation?

Decreasing latency of strong consistency

What is the
latency of the **fetch_timeline()** operation?

Baseline

Read using a quorum of 2/3 replicas

vs.

ICG

1. *Weak*: Read with 1/3 replicas
2. “*Strong*”: Read with quorum of 2/3 replicas

Decreasing latency of strong consistency

What is the latency of the **fetch_timeline()** operation?

Workload A (50:50 read/write)

Workload B (95:5 read/write)

Workload C (read-only)

Baseline

Read using a quorum of 2/3 replicas

vs.


ICG

1. *Weak*: Read with 1/3 replicas
2. “*Strong*”: Read with quorum of 2/3 replicas


Decreasing latency of strong consistency

What is the latency of the **fetch_timeline()** operation?


Workload A (50:50 read/write)



Workload B (95:5 read/write)



Workload C (read-only)



Baseline

Read using a quorum of 2/3 replicas

vs.

ICG


1. *Weak*: Read with 1/3 replicas
2. *“Strong”*: Read with quorum of 2/3 replicas

Decreasing latency of strong consistency


What is the latency of the **fetch_timeline()** operation?

- ★ **Latency decrease by 40%**
- ★ *Throughput drop by 6%*
- ★ *Same consistency model (2/3 replicas)*


Workload A (50:50 read/write)



Workload B (95:5 read/write)



Workload C (read-only)



Baseline

Read using a quorum of 2/3 replicas

vs.


ICG

1. *Weak:* Read with 1/3 replicas
2. *“Strong:”* Read with quorum of 2/3 replicas

Conclusion

The Correctables abstraction enables you to:

- 1. Leverage consistency models incrementally**
- 2. Lower latency of strong consistency**



backup slides

Speculation // Syntactic sugar

```
1 invoke(read(...))  
2   .speculate(speculationFunc[, abortFunc])  
3   .setCallbacks(onFinal = (res) => deliver(res))
```

Listing 3: Generic speculation with Correctables. The square brackets indicate that abortFunc is optional.

Legacy code vs. Correctables


```
1 from pylons import app_globals as g # cache access
2 from r2.lib.db import queries      # backend access

4 def user_messages(user, update = False):
5     key = messages_key(user._id)
6     trees = g.permacache.get(key)
7     if not trees or update:
8         trees = user_messages_nocache(user)
9         g.permacache.set(key, trees) # cache coherence
10    return trees
11 def user_messages_nocache(user):
12     # Just like user_messages, but avoiding the cache...
```

Listing 1: Different consistency guarantees in Reddit [13], as an example of tight coupling between applications and storage. Developers must manually handle the cache and the backend.

```
1 def user_messages(user, strong = False):
2     key = messages_key(user._id)
3     # coherence handled by invoke* functions in bindings
4     if strong: return invokeStrong(get(key))
5     else: return invokeWeak(get(key))
```

Listing 2: Reddit code rewritten using Correctables.



Legacy code vs. Correctables

```
1 from pylons import app_globals as g # cache access
2 from r2.lib.db import queries      # backend access

4 def user_messages(user, update = False):
5     key = messages_key(user._id)
6     trees = g.permacache.get(key)
7     if not trees or update:
8         trees = user_messages_nocache(user)
9         g.permacache.set(key, trees) # cache coherence
10    return trees
11 def user_messages_nocache(user):
12     # Just like user_messages, but avoiding the cache...
```

Listing 1: Different consistency guarantees in Reddit [13], as an example of tight coupling between applications and storage. Developers must manually handle the cache and the backend.

```
1 invoke(getLatestNews()).setCallbacks(
2     onUpdate = (items) => refreshDisplay(items))
```

Listing 6: Progressive display of news items using Correctables. The refreshDisplay function triggers with every update on the news items.

```
1 def user_messages(user, strong = False):
2     key = messages_key(user._id)
3     # coherence handled by invoke* functions in bindings
4     if strong: return invokeStrong(get(key))
5     else: return invokeWeak(get(key))
```


Listing 2: Reddit code rewritten using Correctables.

Correctable

↑ semantics


↓ execution details

Overheads




- Cassandra
- YCSB workload, various configurations
- Client in Ireland
- Replicas in Virginia, Frankfurt, and Ireland

Overheads



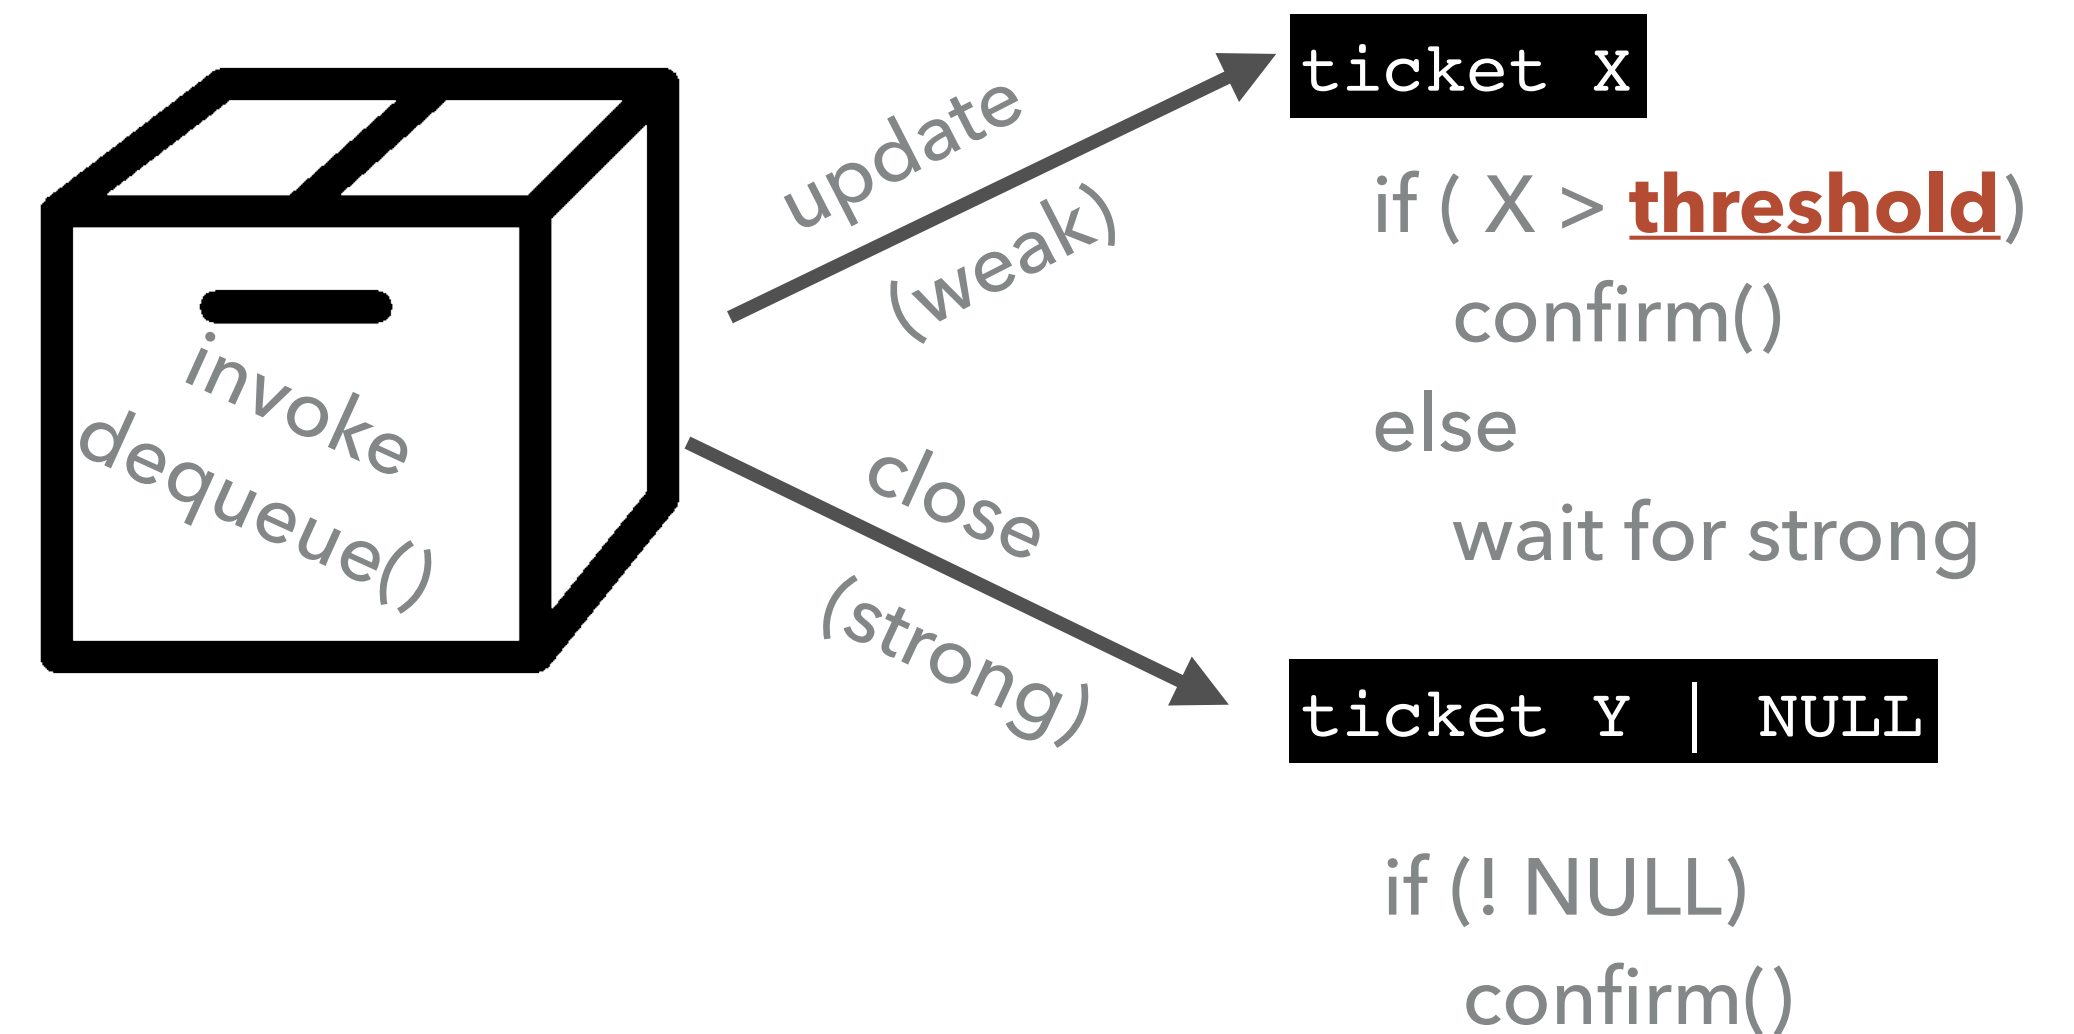

- ZooKeeper queue implementation
- Wasteful implementation (by default)
- We were able to improve — negative overhead

- Cassandra
- YCSB workload, various configurations
- Client in Ireland
- Replicas in Virginia, Frankfurt, and Ireland



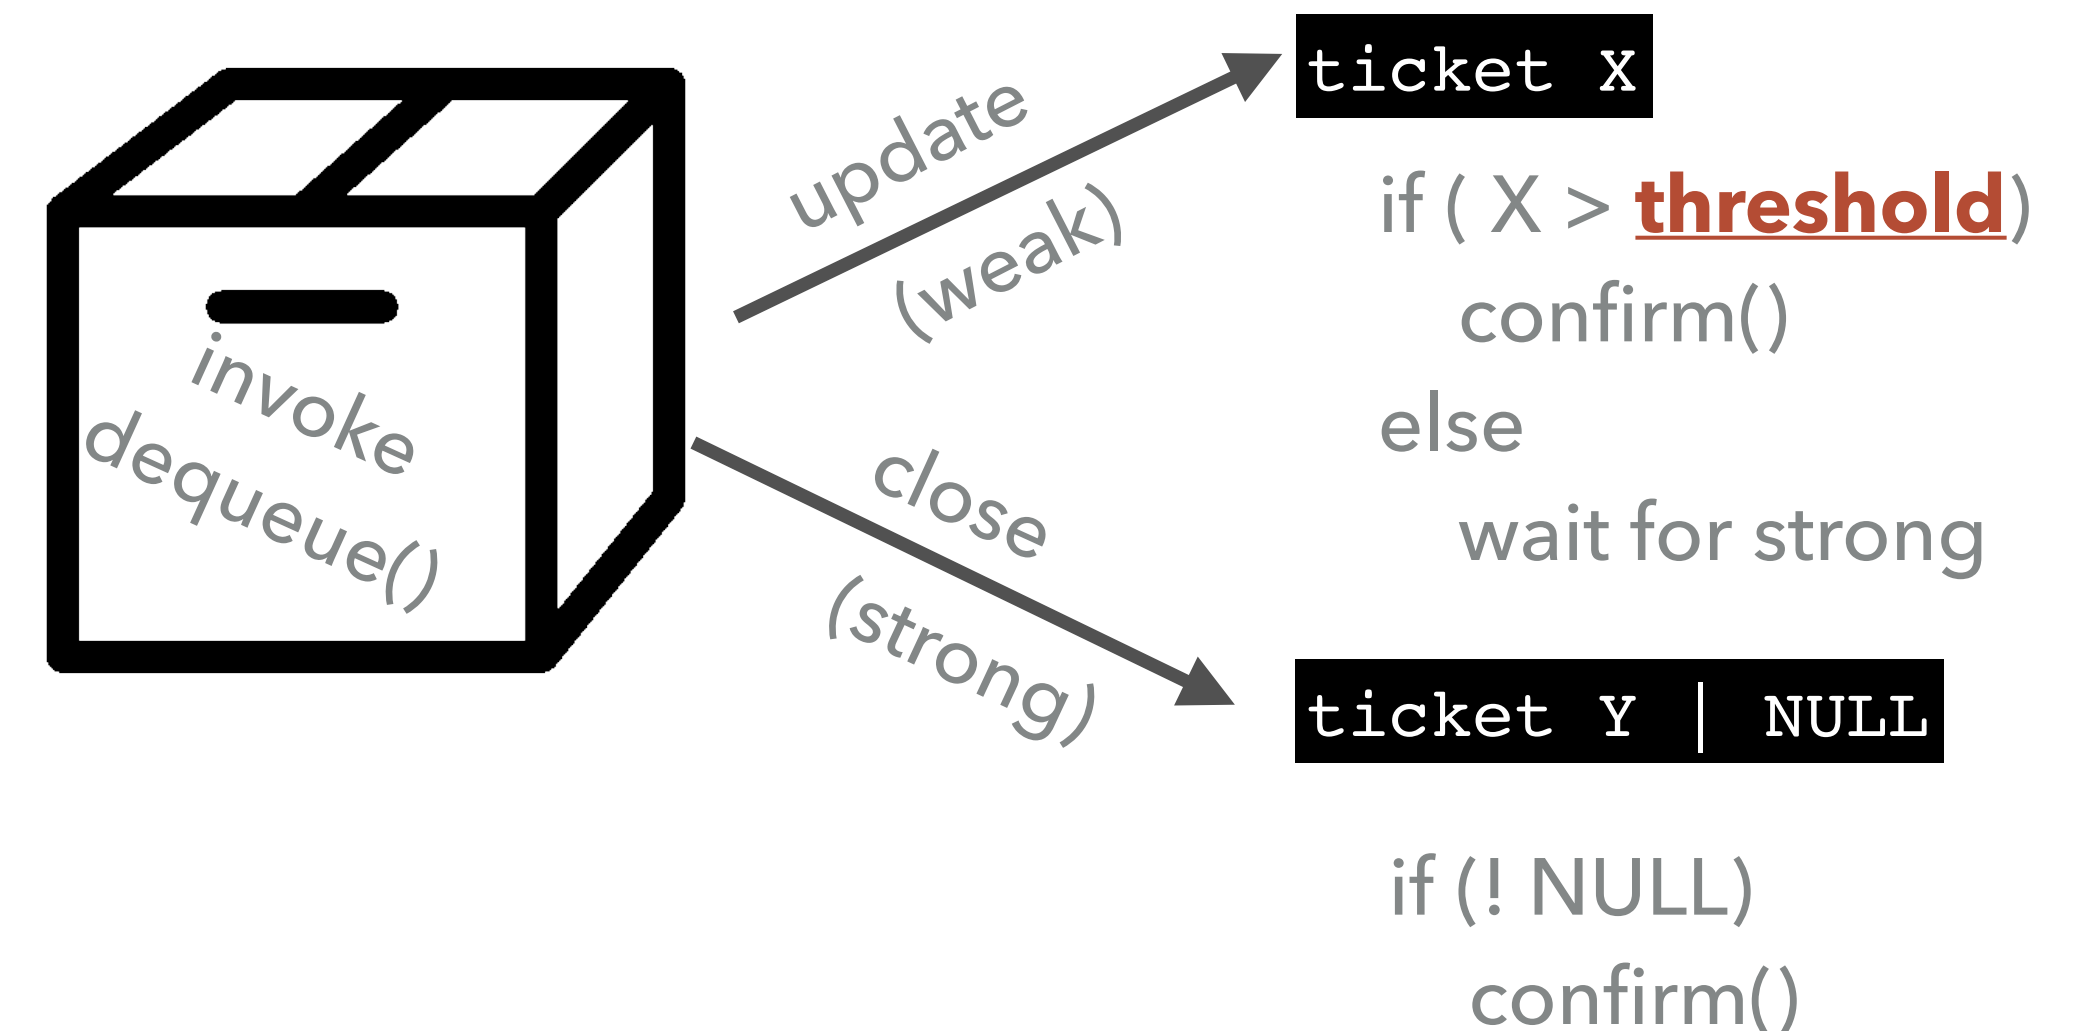

Exploiting application semantics

- Ticket selling application
 - Implemented through a ZooKeeper queue
 - Buy ticket = dequeue operation




Exploiting application semantics

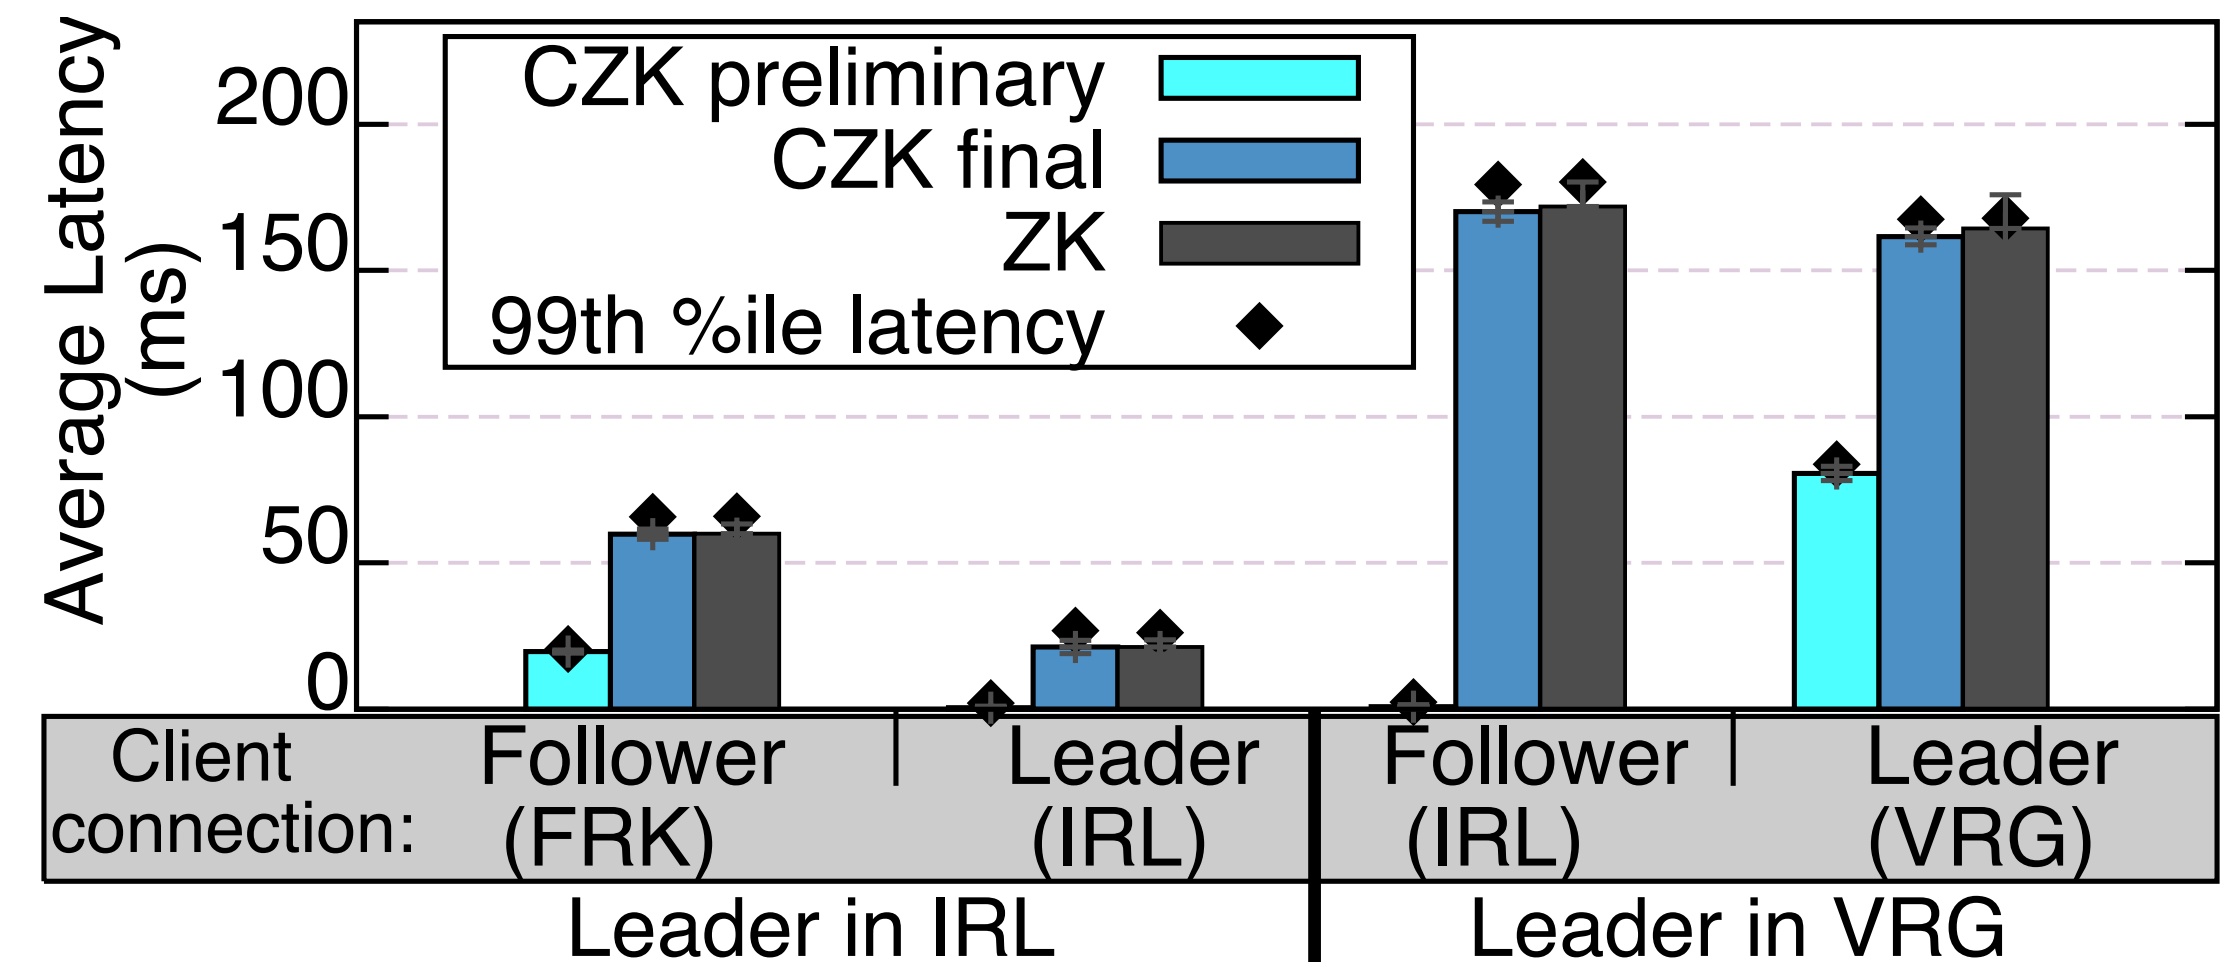

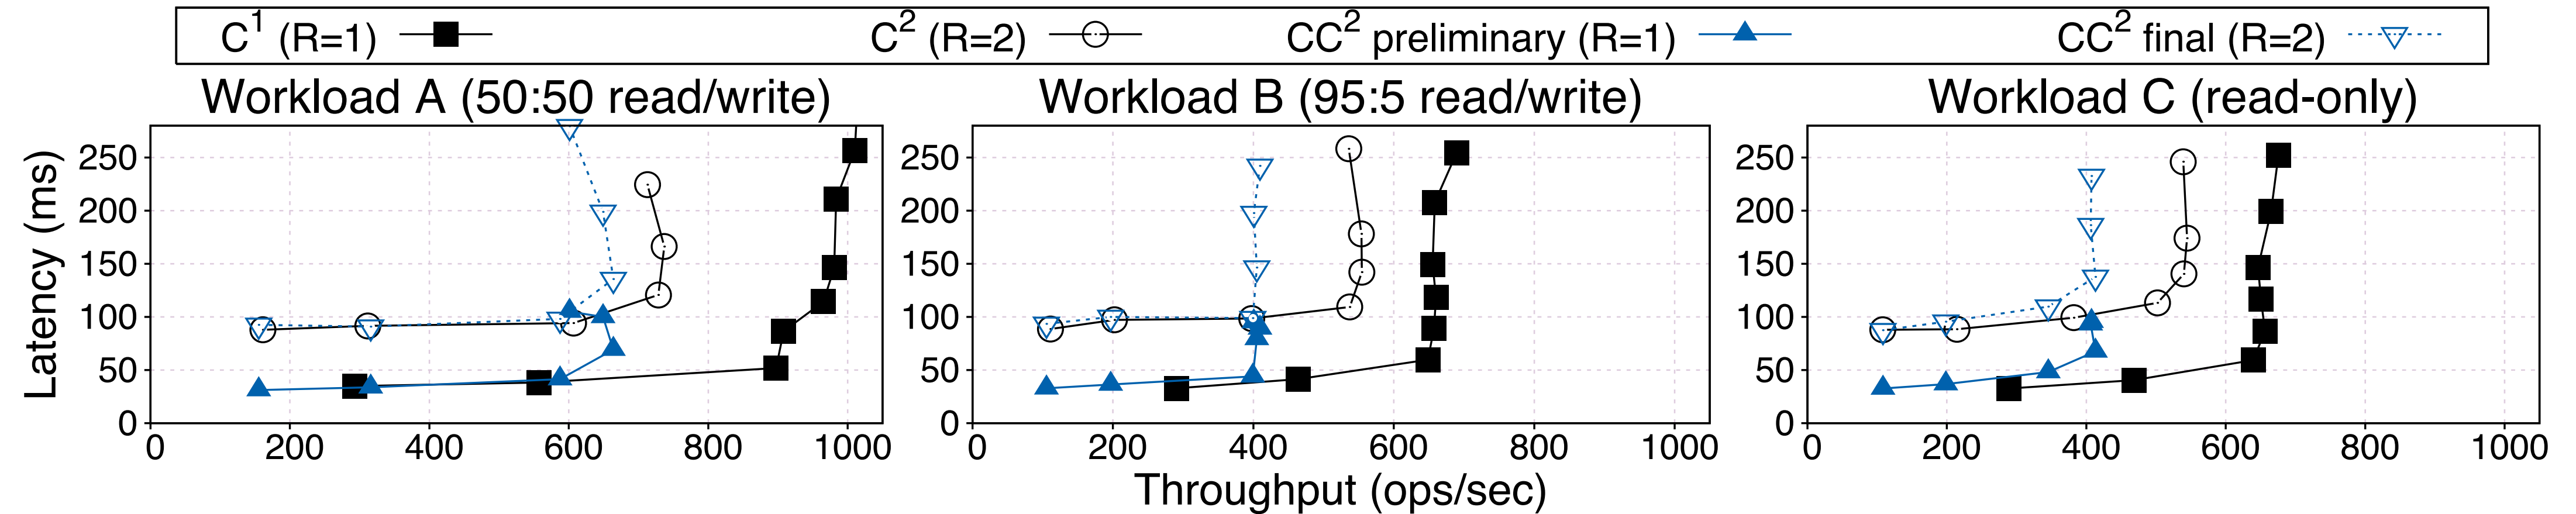
- Ticket selling application
 - Implemented through a ZooKeeper queue
 - Buy ticket = dequeue operation




Divergence between weak and strong consistency



Latency gaps between consistency models



Efficiency of Multiple Responses



Efficiency of Multiple Responses

SOCIAL MEDIA APPLICATION

read
timeline

value¹

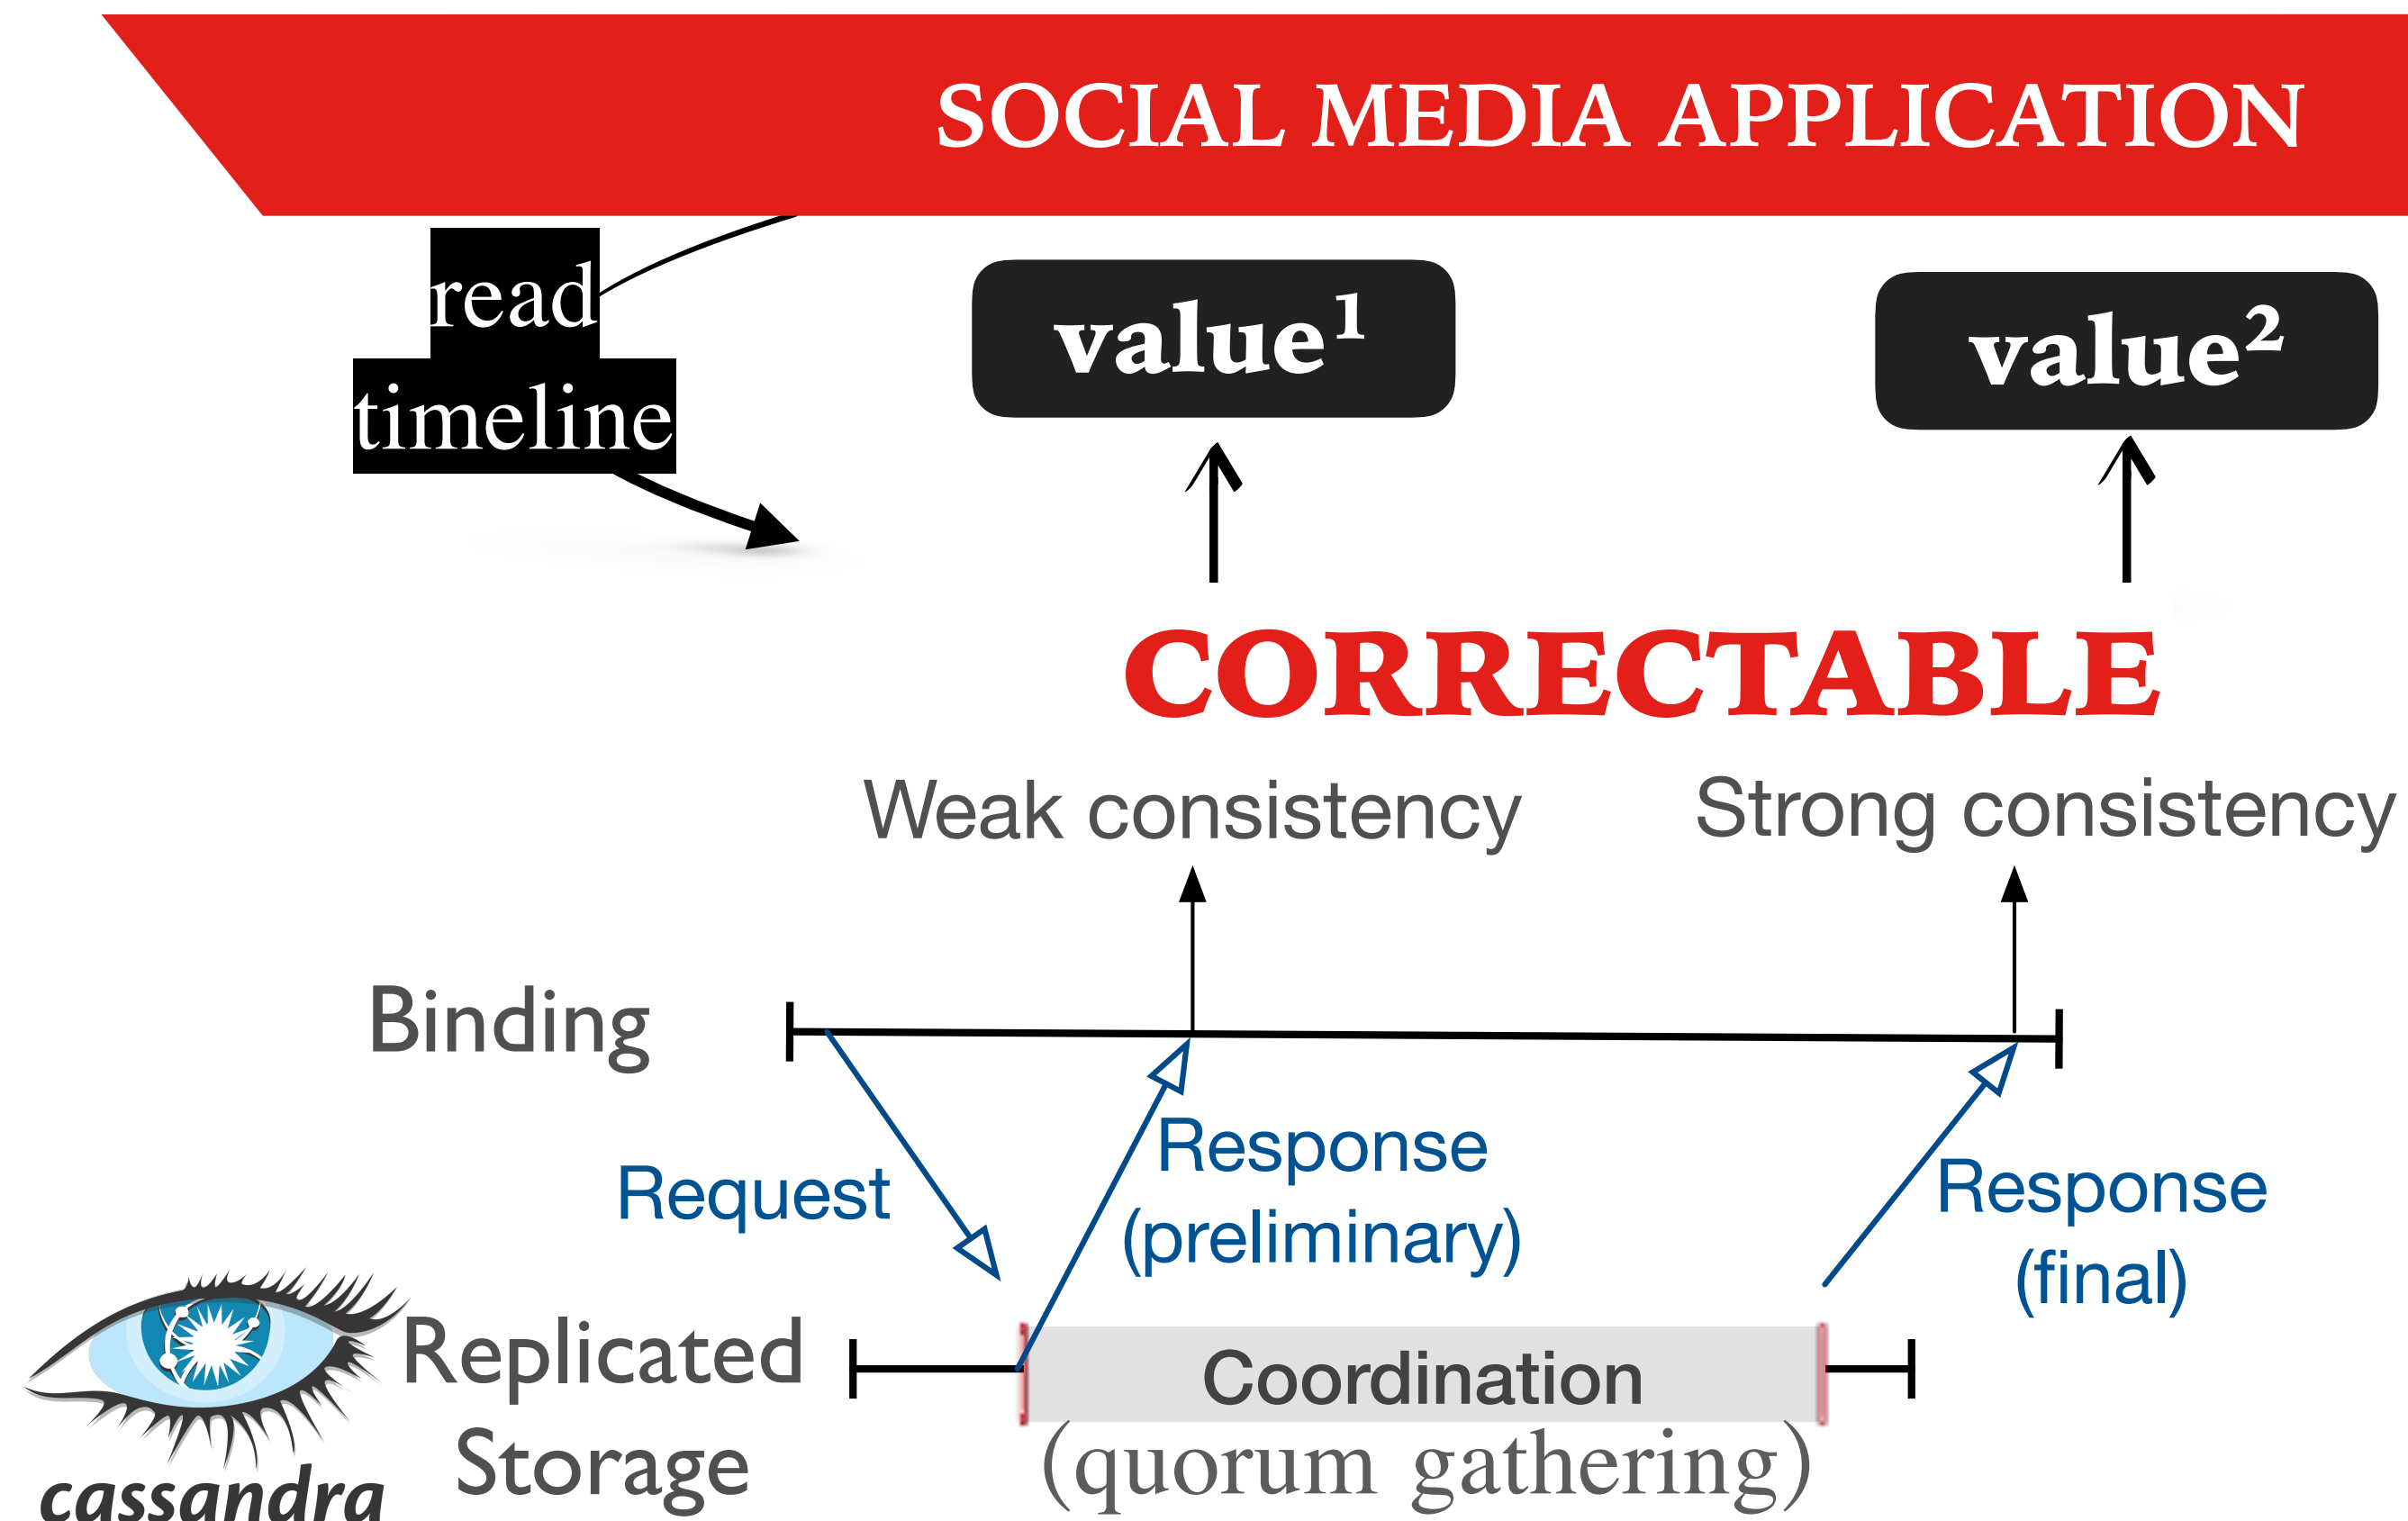
value²

CORRECTABLE

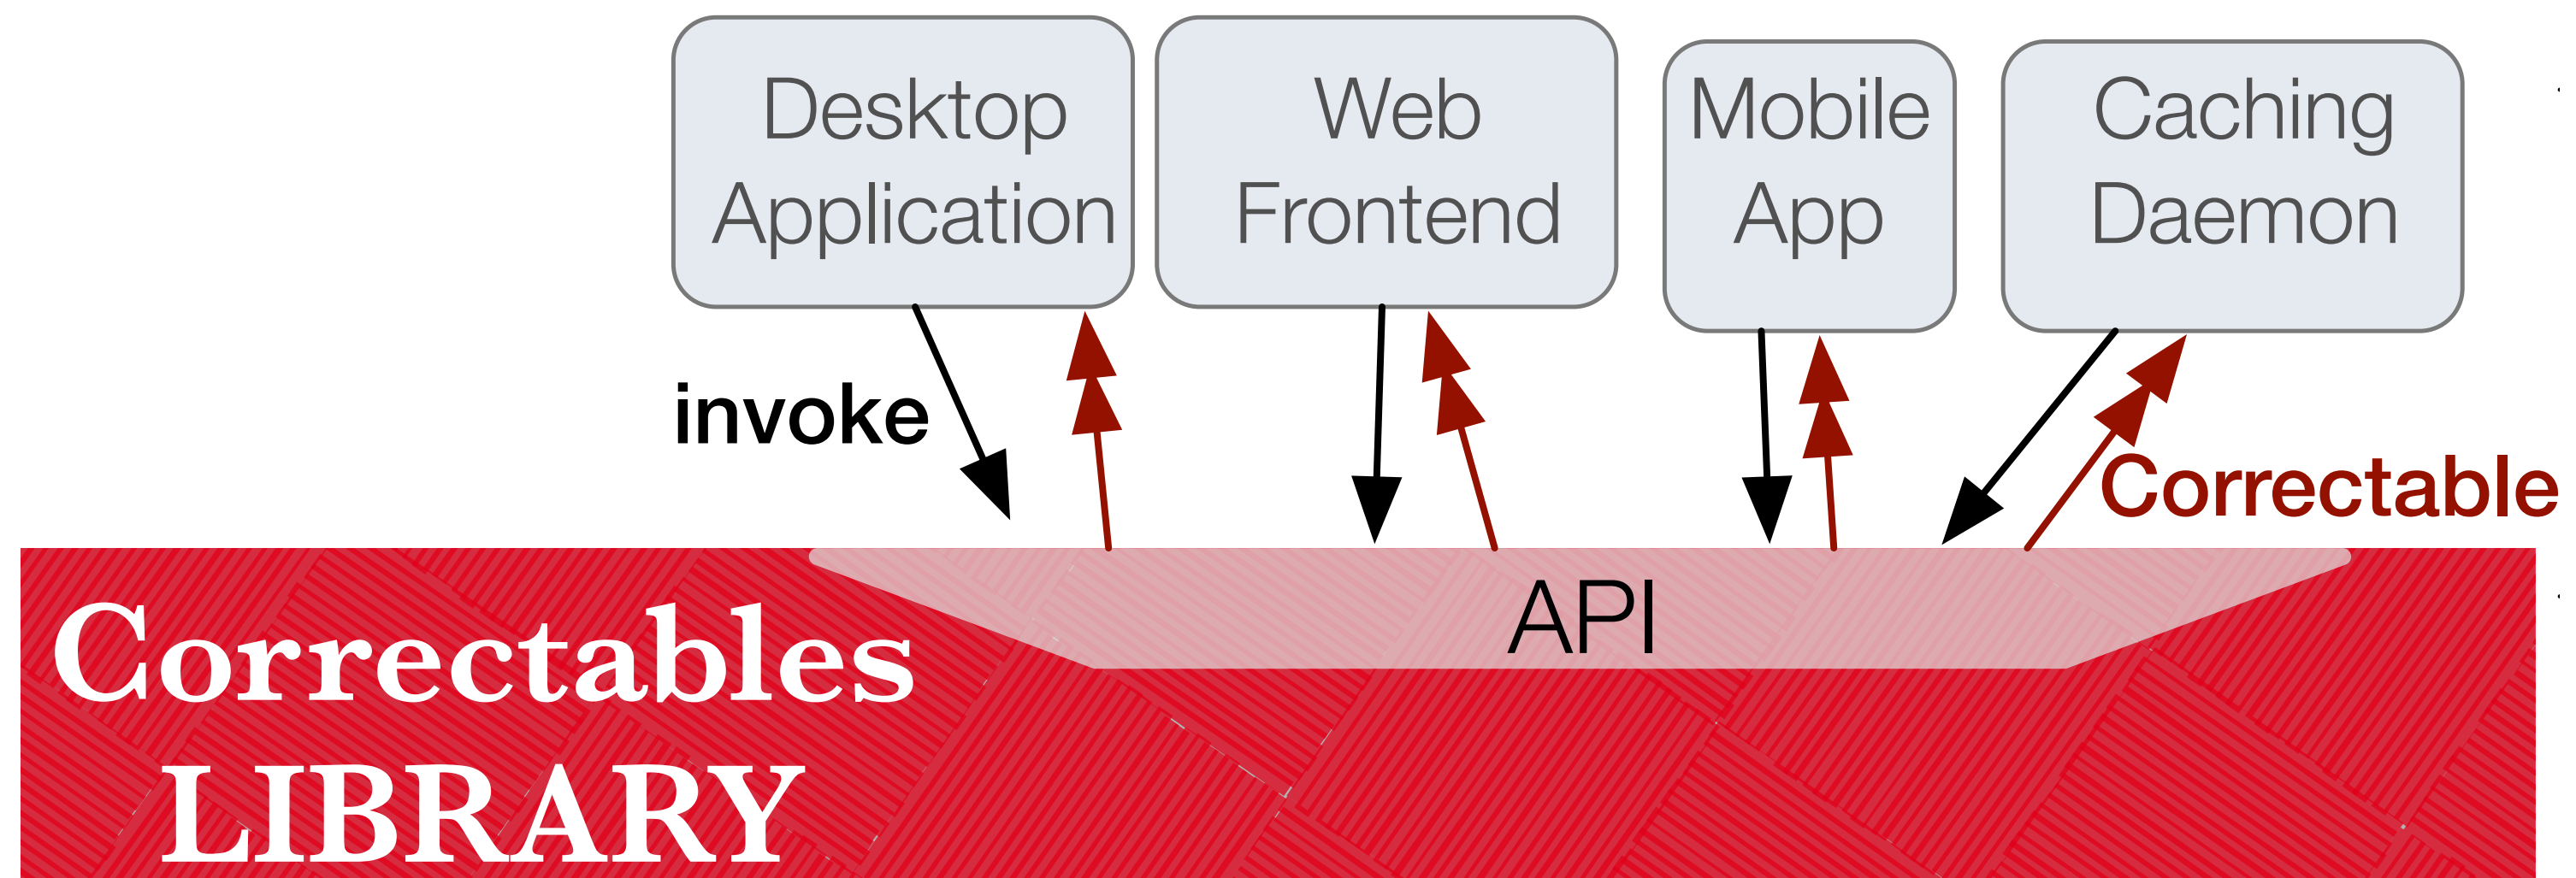
Binding



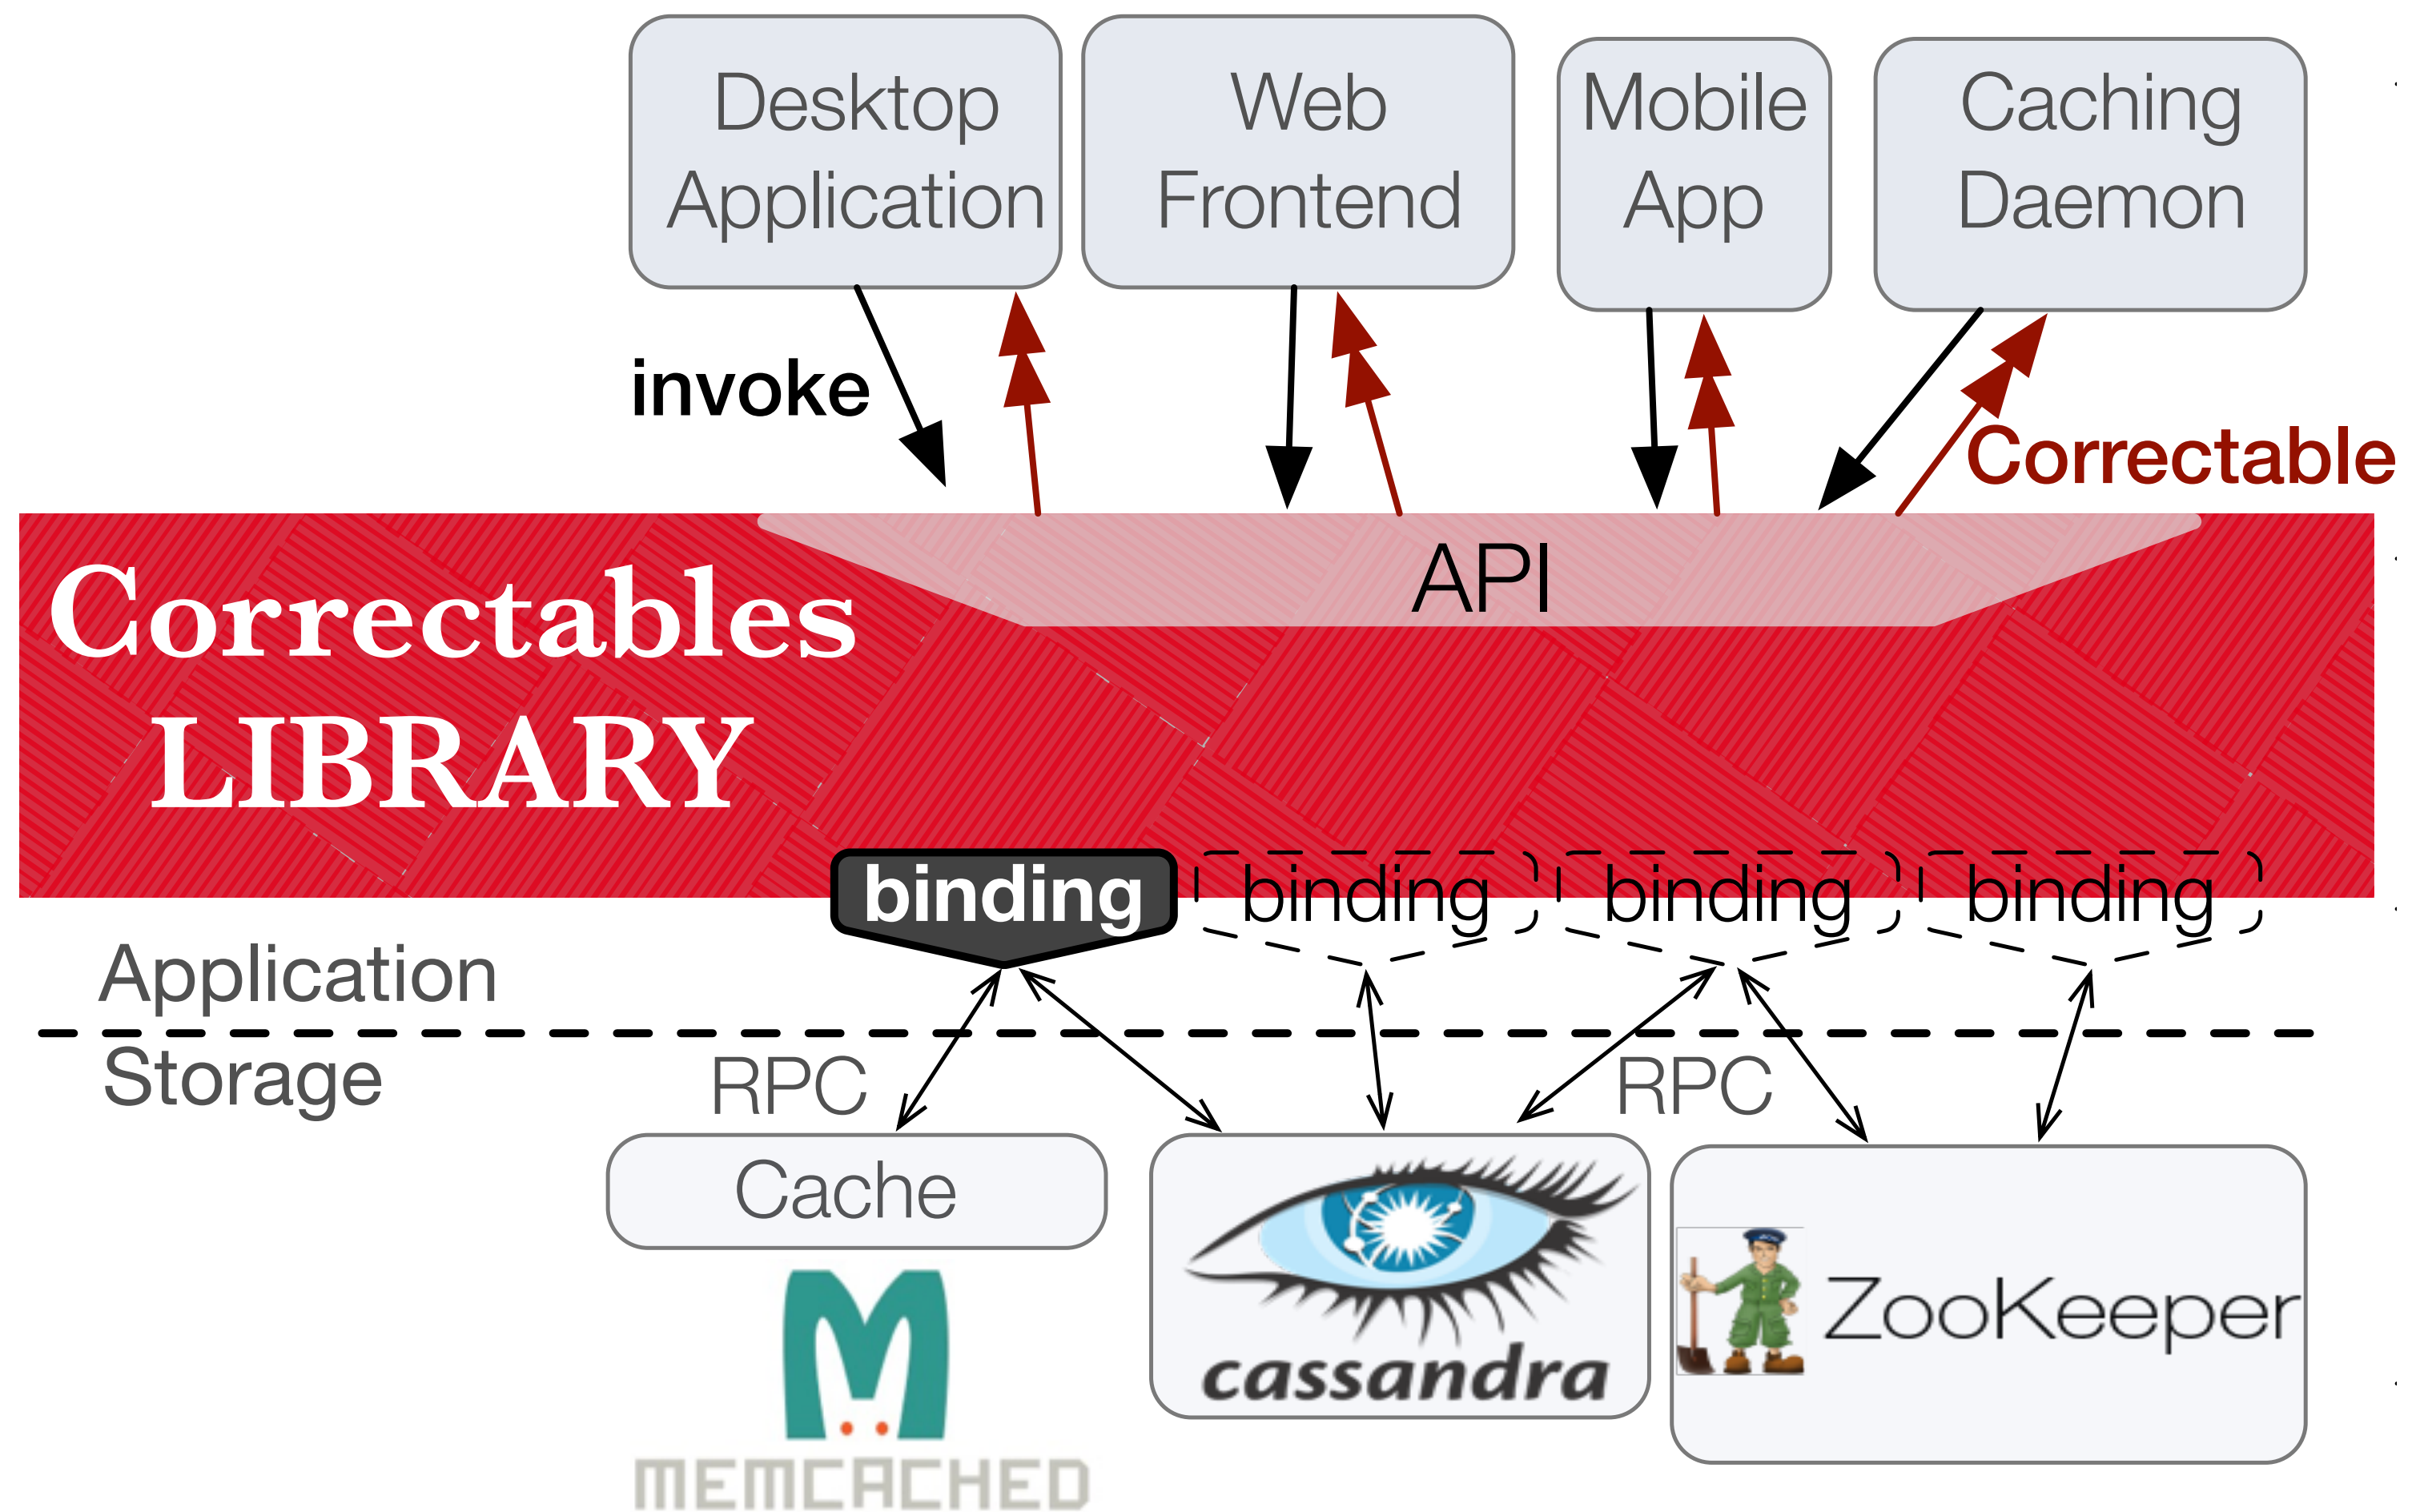
Efficiency of Multiple Responses



Correctables / **Library**



Correctables / **Library**




- ★ Infrastructure services
- ★ Stock tickers
- ★ Trading applications



- ★ E-mail
- ★ Calendar
- ★ Social networks
- ★ Online shopping
- ★ Ad serving
- ★ News reading
- ★ Collaborative editing

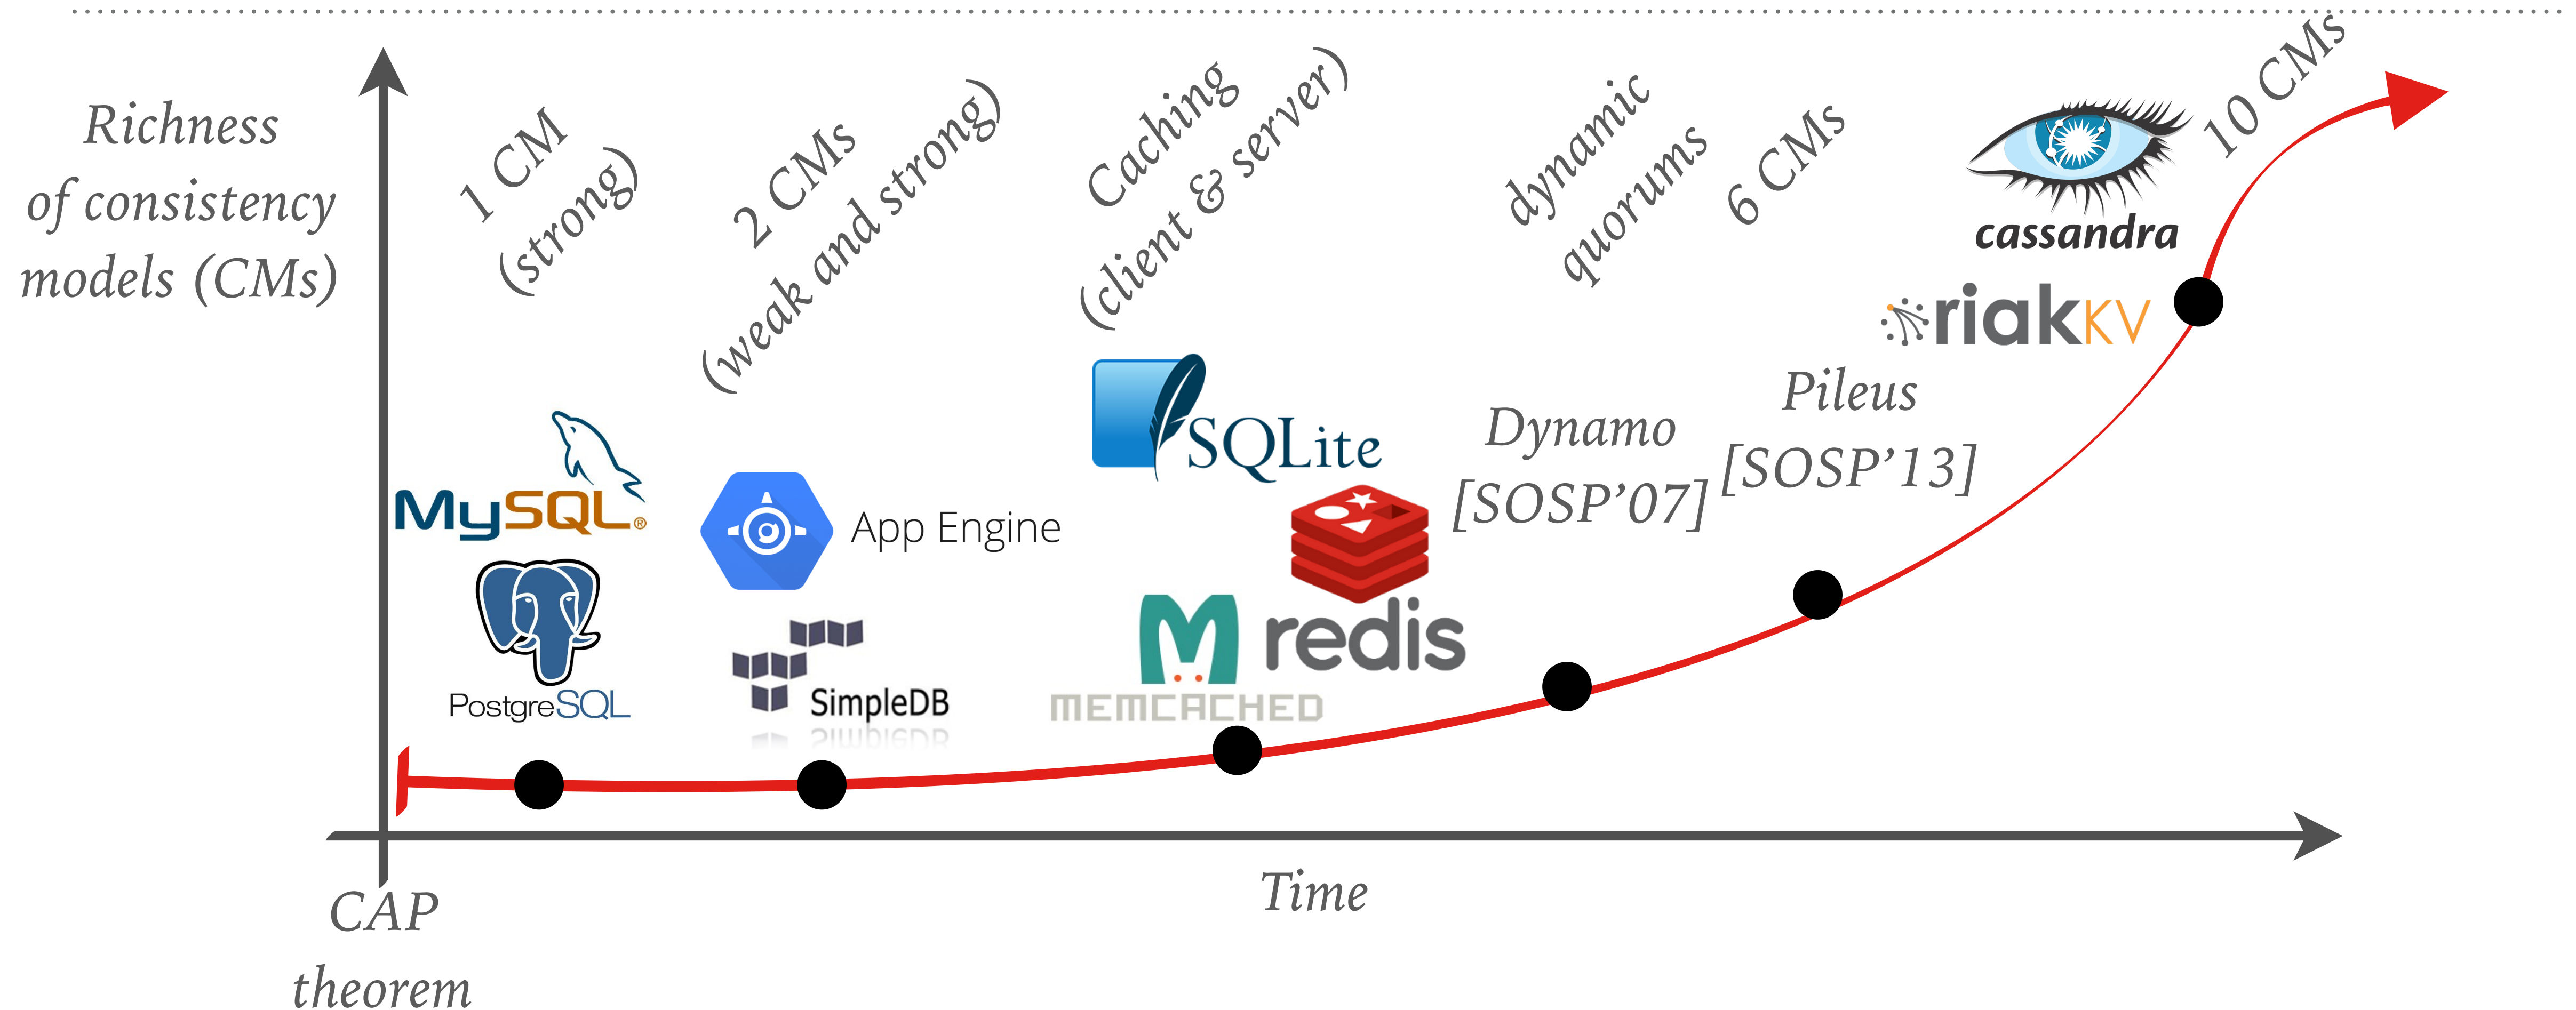
Performance is a second-order concern



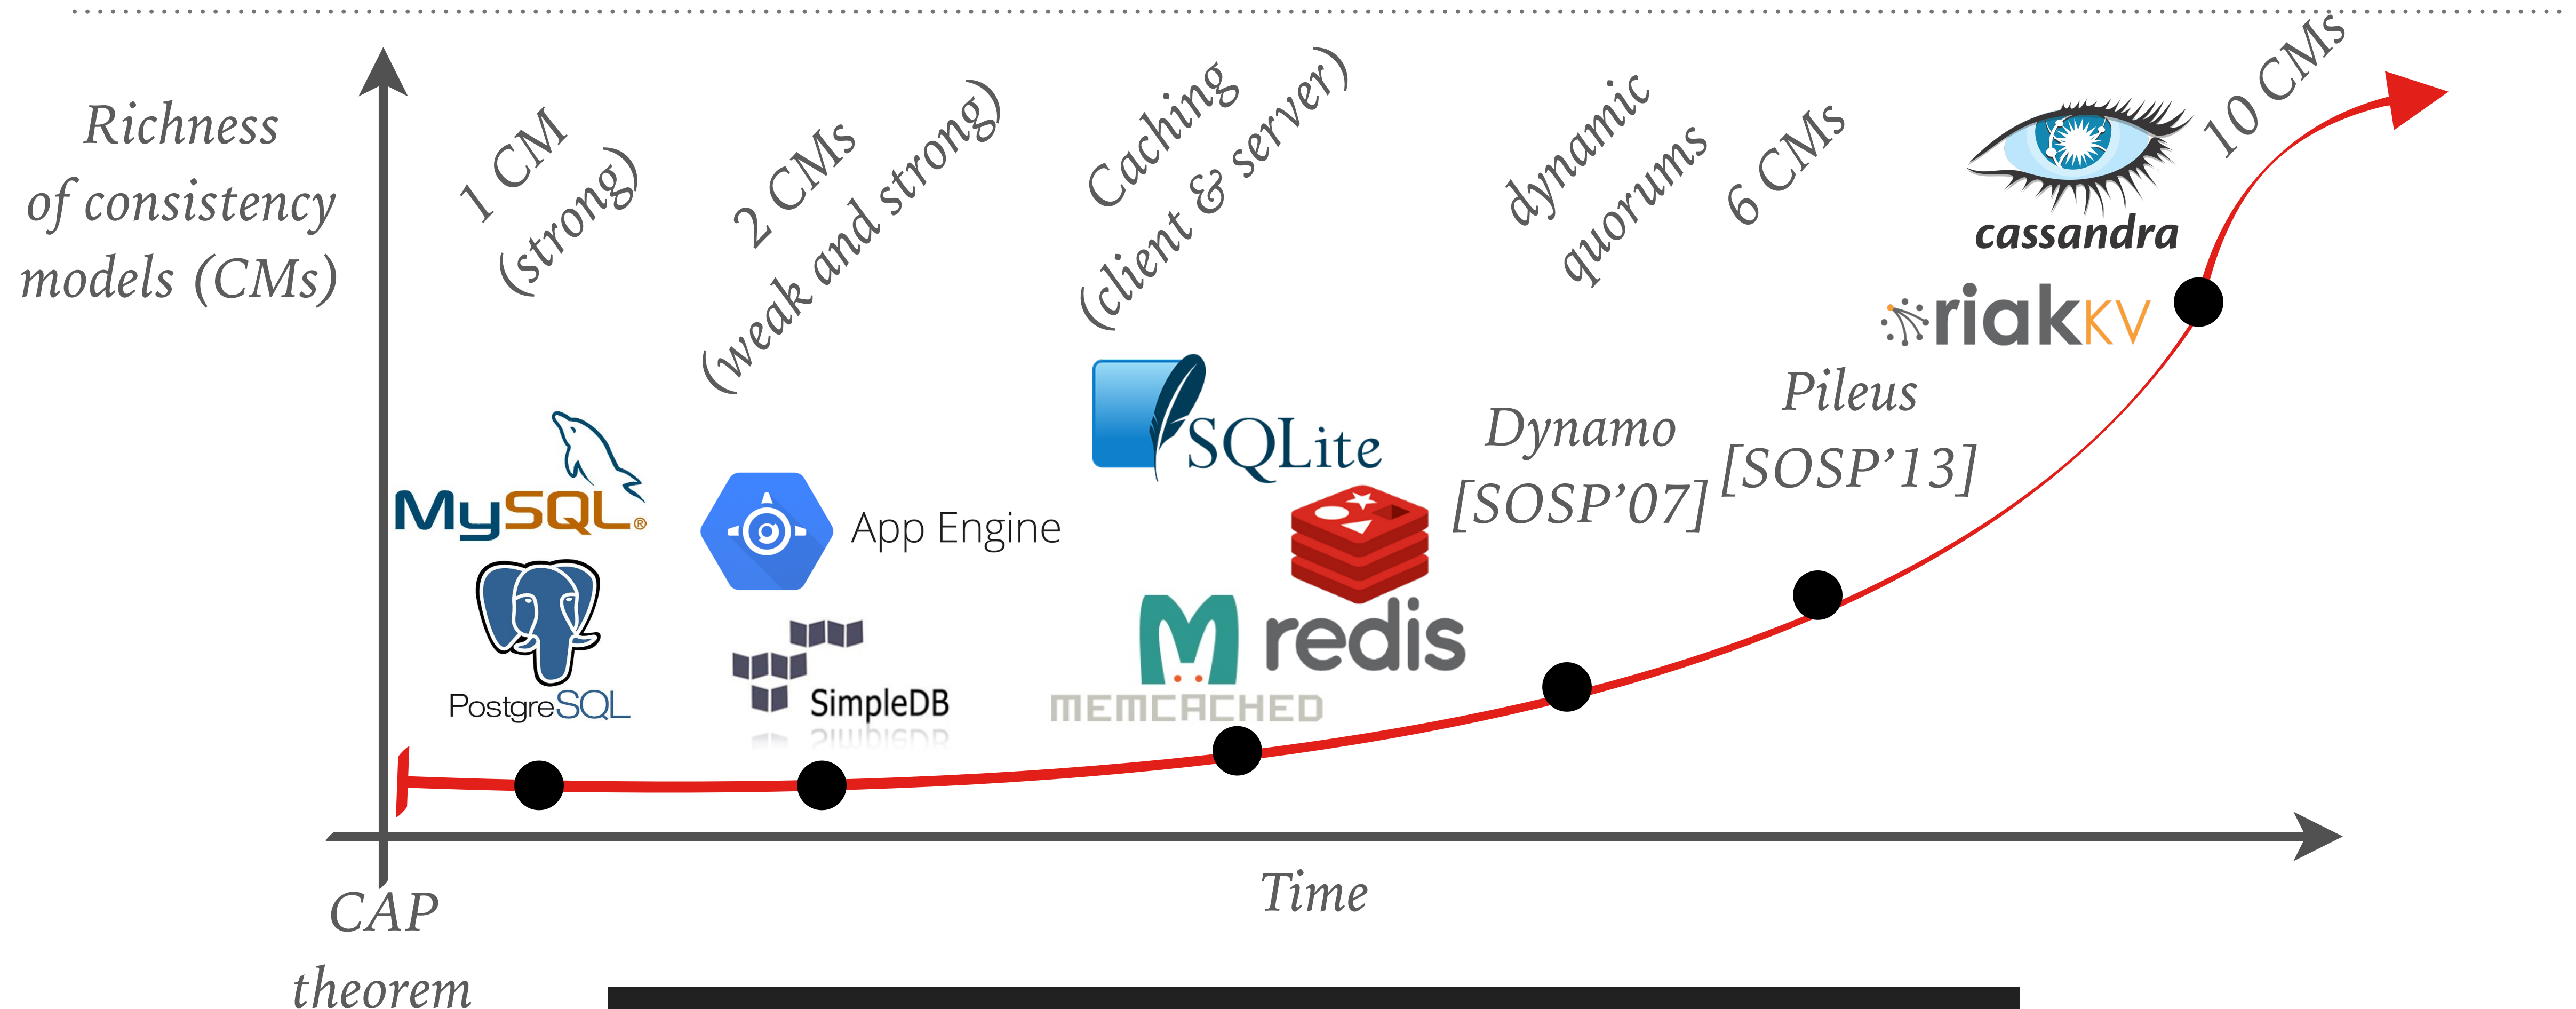
- ★ Computation on static content
- ★ Cold data analysis
- ★ Disconnected operations in mobile applications

Semantics allow
to bypass the need for strong
consistency

Replicated storage systems



Replicated storage systems



No single consistency model is ideal
→ expose multiple choices